

Lecture 3

Processes

Prof. Yinqian Zhang

Spring 2023


Outline

- Process and system calls
- Process creation
- Kernel view of processes
- Kernel view of `fork()`, `exec()`, and `wait()`
- More about processes
- Threads

Process and System Calls


What Is a Process

- Process is a program in execution
- A program is a file on the disk
 - Code and static data
- A process is loaded by the OS
 - Code and static data are loaded from the program
 - Heap and stack are created by the OS



What Is a Process (Cont'd)

- A process is an abstraction of machine states
 - Memory: address space
 - Register:
 - Program Counter (PC) or Instruction Pointer
 - Stack pointer
 - frame pointer
 - I/O: all files opened by the process



Process Identification

- How can we distinguish processes from one to another?
 - Each process is given a unique ID number, and is called the process ID, or the PID.
 - The system call, `getpid()`, prints the PID of the calling process.

```
// compile to getpid
#include <stdio.h>    // printf()
#include <unistd.h>   // getpid()

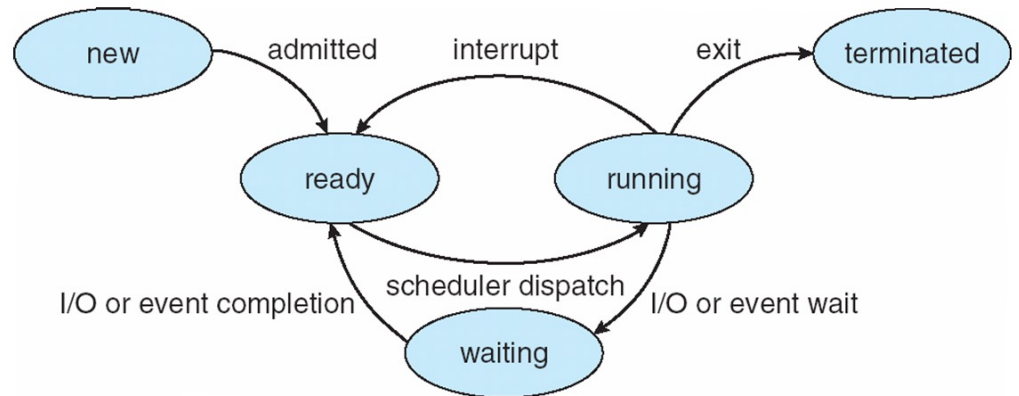
int main(void) {
    printf("My PID is %d\n", getpid() );
}
```

```
$ ./getpid
My PID is 1234
$ ./getpid
My PID is 1235
$ ./getpid
My PID is 1237
```

Process Life Cycle

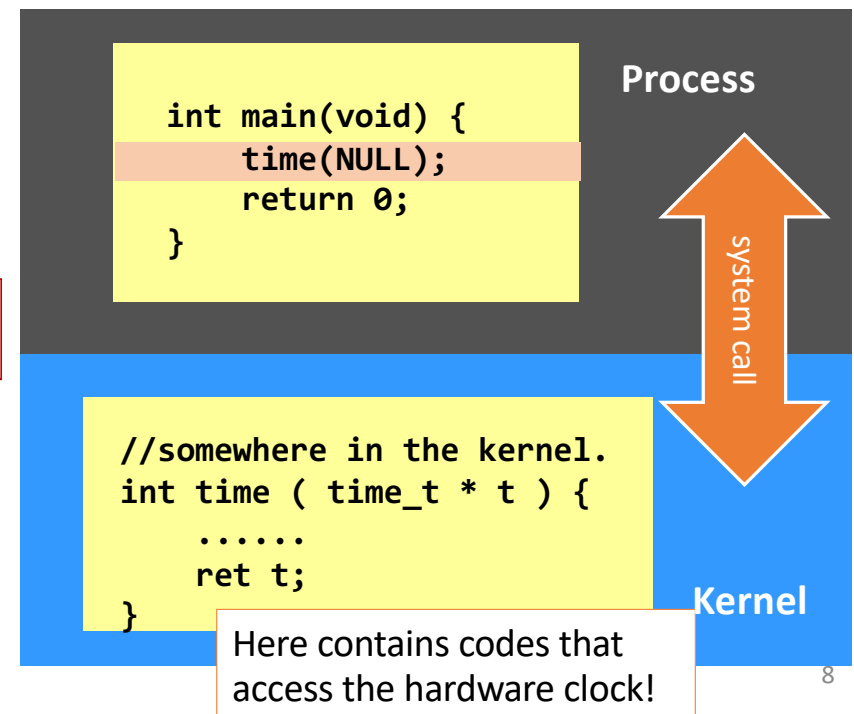
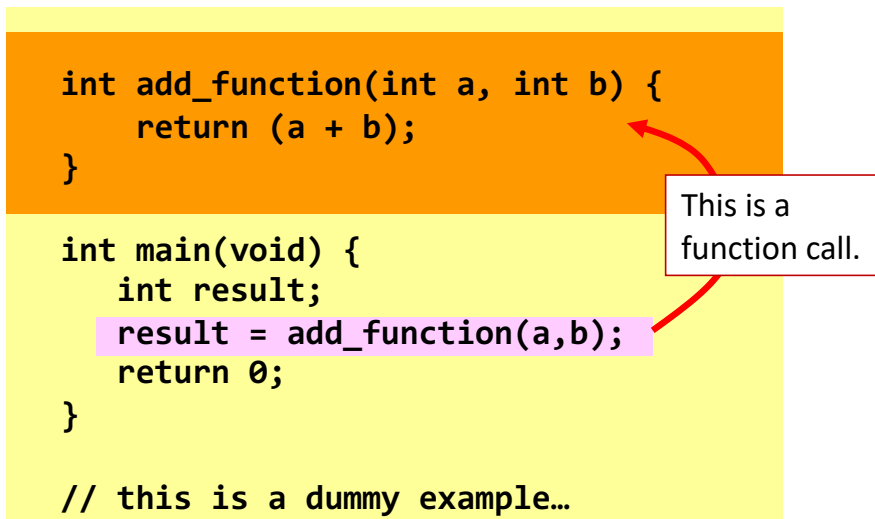
```
int main(void) {  
    int x = 1;  
    getchar();  
    return x;  
}
```

Process State



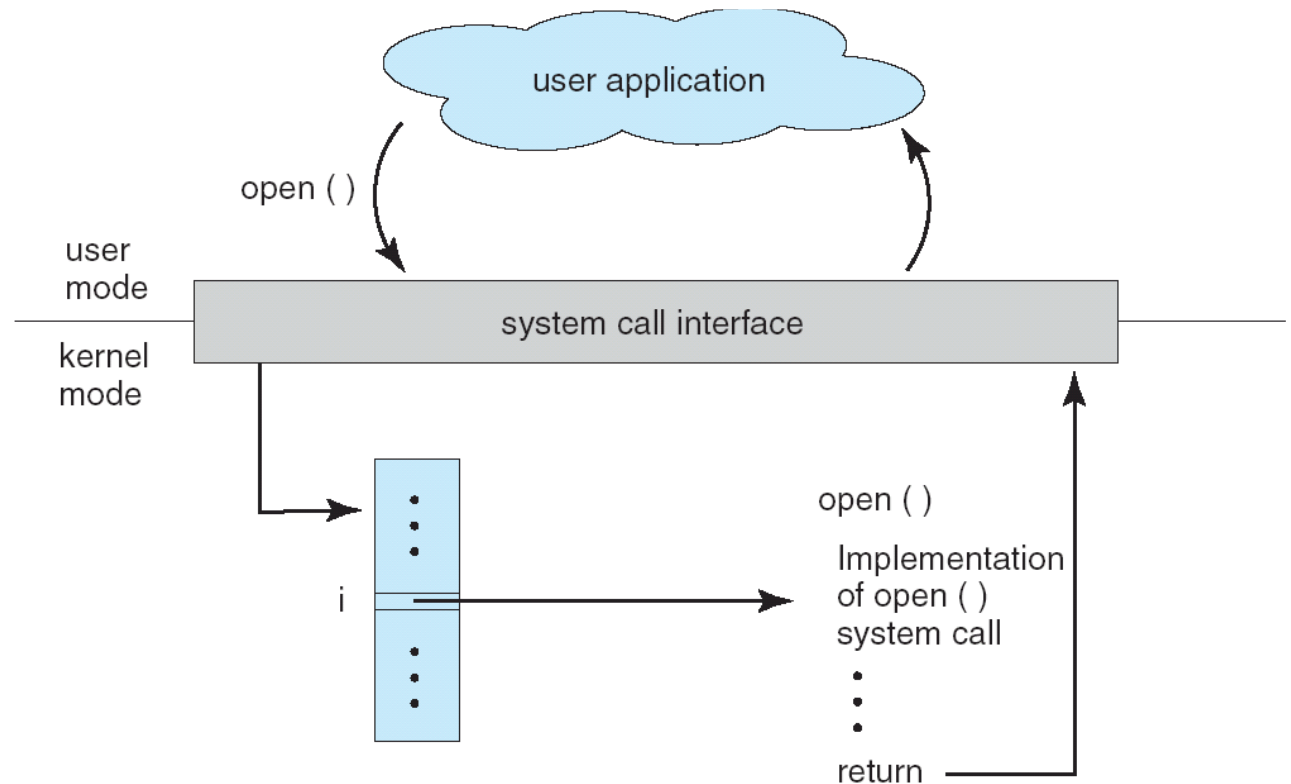
System Call: Process-Kernel Interaction

- System call is a function call.
 - exposed by the **kernel**.
 - abstraction of kernel operations.



System Call: Call by Number

- System call is different from function call
- System call is a call by number



System Call: Call by Number

- User-mode code from xv6-riscv

```
int main(void) {  
    .....  
    int fd = open("copyin1", O_CREATE|O_WRONLY);  
    .....  
    return 0;  
}
```

```
/* kernel/syscall.h */
```

```
#define SYS_open 15
```

```
/* user/usys.S */
```

```
.global open
```

```
open:
```

```
    li a7, SYS_open
```

```
    ecall
```

```
    ret
```

System Call: Call by Number

- Kernel code from xv6-riscv

```
/* kernel/syscall.h */
```

```
#define SYS_open 15
```

```
/* kernel/file.c */
```

```
uint64 sys_open(void) {  
    .....  
    return fd;  
}
```

```
/* kernel/syscall.c */
```

```
static uint64 (*syscalls[])(void) = {  
    .....  
    [SYS_open] sys_open,  
    .....  
}
```

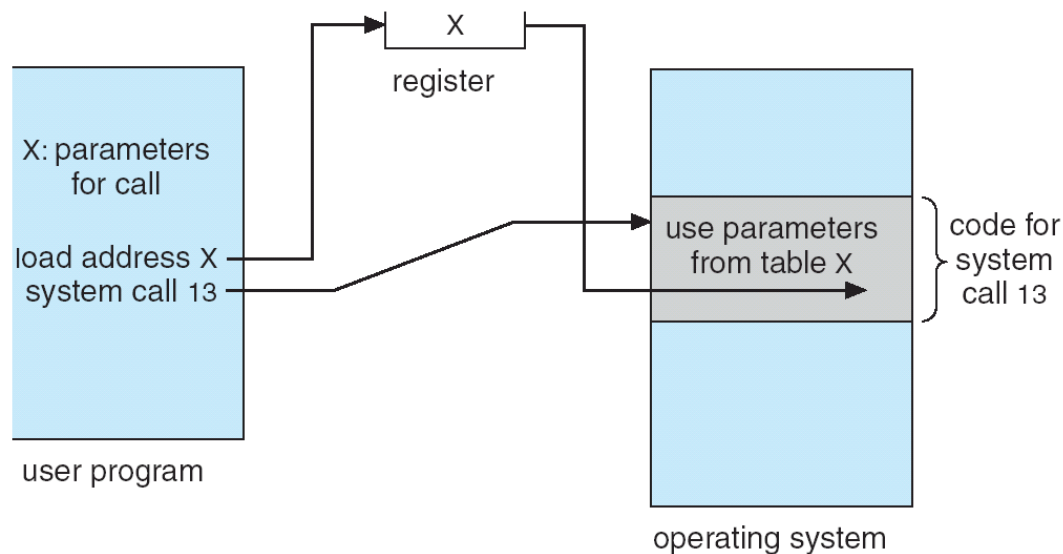
```
void syscall(void) (  
    struct proc *p = myproc();  
    num = p->trapframe->a7;  
    p->trapframe->a0 = syscalls[num]();  
}
```

System Call: Parameter Passing

- Often, more information is required than the index of desired system call
 - Exact type and amount of information vary according to OS and call
- Three general methods used to pass parameters to the OS
 - **Registers:** pass the parameters in registers
 - In some cases, may be more parameters than registers
 - x86 and risc-v take this approach
 - **Blocks:** Parameters stored in a memory block and address of the block passed as a parameter in a register
 - **Stack:** Parameters placed, or pushed, onto the stack by the program and popped off the stack by the operating system
 - Block and stack methods do not limit the number or length of parameters being passed

System Call: Parameter Passing


- Example: parameter passing via blocks



System Call v.s. Library API Call

- Most operating systems provide standard C library to provide library API calls
 - A layer of indirection for system calls

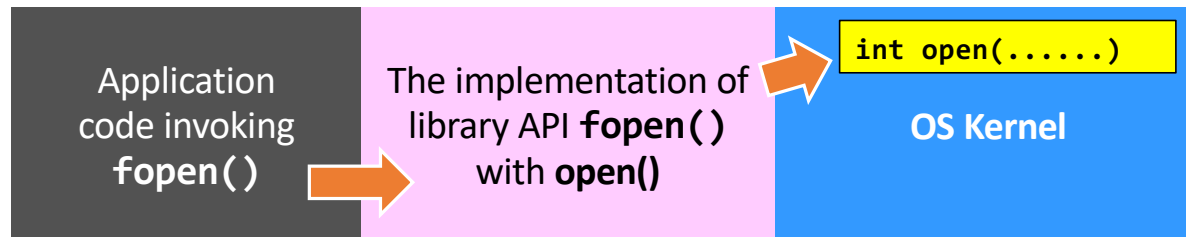
Name	System call?
printf() & scanf()	No
malloc() & free()	No
fopen() & fclose()	No
mkdir() & rmdir()	Yes
chown() & chmod()	Yes



System Call v.s. Library API Call

- Take `fopen()` as an example.
 - `fopen()` invokes the system call `open()`.
 - `open()` is too primitive and is not programmer-friendly!

Library call	<code>fopen("hello.txt", "w");</code>
System call	<code>open("hello.txt", O_WRONLY O_CREAT O_TRUNC, 0666);</code>



Process Creation

Process Creation


- Parent process create children processes, which, in turn create other processes, forming a tree of processes
- Generally, process identified and managed via a process identifier (pid)
- Resource sharing
 - Parent and children share all resources
 - Children share subset of parent's resources
 - Parent and child share no resources
- Execution
 - Parent and children execute concurrently
 - Parent waits until children terminate

Process Creation (Cont'd)

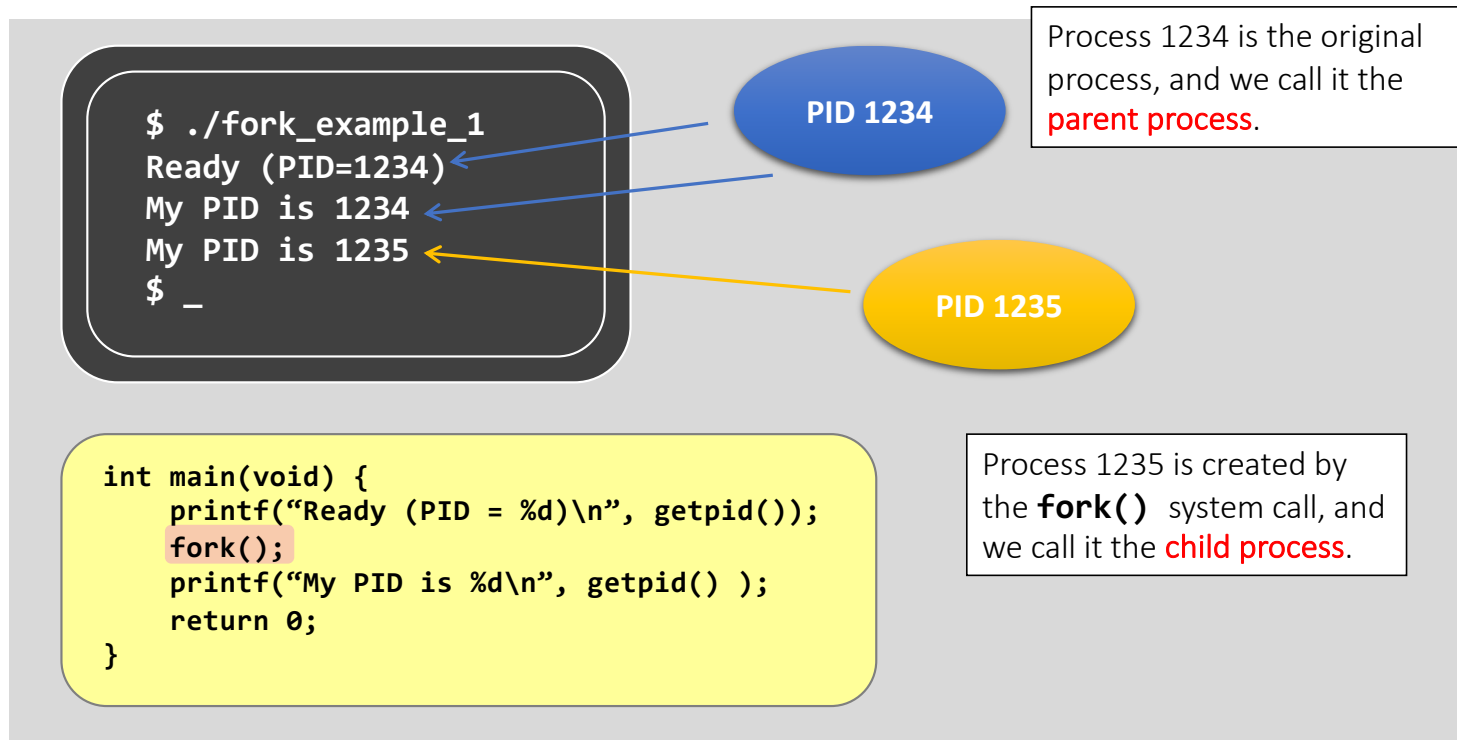
- Address space
 - Child duplicate of parent
 - Child has a program loaded into it
- UNIX examples
 - fork system call creates new process
 - exec system call used after a fork to replace the process' memory space with a new program

Process Creation (Cont'd)

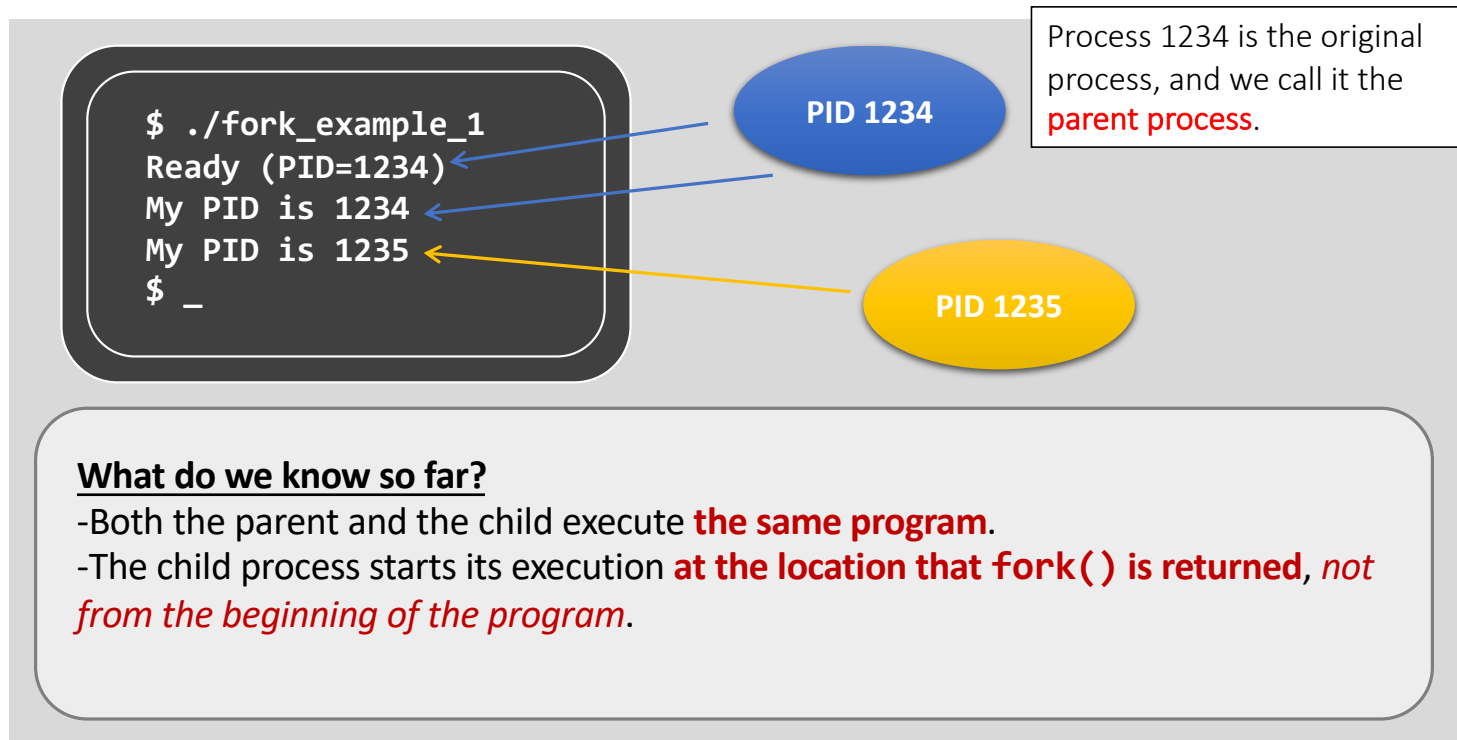
- A tree of processes in Linux




Creating Processes with fork() System Call



Creating Processes with fork() System Call



Creating Processes with fork() System Call




```
1 int main(void) {  
2     int result;  
3     printf("before fork ...\n");  
4     result = fork();  
5     printf("result = %d.\n", result);  
6  
7     if(result == 0) {  
8         printf("I'm the child.\n");  
9         printf("My PID is %d\n", getpid());  
10    }  
11    else {  
12        printf("I'm the parent.\n");  
13        printf("My PID is %d\n", getpid());  
14    }  
15  
16    printf("program terminated.\n");  
17 }
```

```
$ ./fork_example_2  
before fork ...
```

PID 1234

Creating Processes with fork() System Call



```
1 int main(void) {
2     int result;
3     printf("before fork ...\n");
4     result = fork();
5     printf("result = %d.\n", result);
6
7     if(result == 0) {
8         printf("I'm the child.\n");
9         printf("My PID is %d\n", getpid());
10    }
11    else {
12        printf("I'm the parent.\n");
13        printf("My PID is %d\n", getpid());
14    }
15
16    printf("program terminated.\n");
17 }
```

```
$ ./fork_example_2
before fork ...
```

Important

- Both parent and child need to return from fork().
- CPU scheduler decides which to run first.

PID 1234

fork()

PID 1235

Creating Processes with fork() System Call

```
1 int main(void) {
2     int result;
3     printf("before fork ...\n");
4     result = fork();
5     printf("result = %d.\n", result);
6
7     if(result == 0) {
8         printf("I'm the child.\n");
9         printf("My PID is %d\n", getpid());
10    }
11    else {
12        printf("I'm the parent.\n");
13        printf("My PID is %d\n", getpid());
14    }
15
16    printf("program terminated.\n");
17 }
```

```
$ ./fork_example_2
before fork ...
result = 1235
```

Important

For parent, the return value of **fork()** is the PID of the created child.

PID 1234
(running)

PID 1235
(waiting)

Creating Processes with fork() System Call


```
1 int main(void) {  
2     int result;  
3     printf("before fork ...\n");  
4     result = fork();  
5     printf("result = %d.\n", result);  
6  
7     if(result == 0) {  
8         printf("I'm the child.\n");  
9         printf("My PID is %d\n", getpid());  
10    }  
11    else {  
12        printf("I'm the parent.\n");  
13        printf("My PID is %d\n", getpid());  
14    }  
15  
16    printf("program terminated.\n");  
17 }
```

```
$ ./fork_example_2  
before fork ...  
result = 1235  
I'm the parent.  
My PID is 1234
```


PID 1234
(running)

PID 1235
(waiting)

Creating Processes with fork() System Call



```
1 int main(void) {
2     int result;
3     printf("before fork ...\n");
4     result = fork();
5     printf("result = %d.\n", result);
6
7     if(result == 0) {
8         printf("I'm the child.\n");
9         printf("My PID is %d\n", getpid());
10    }
11    else {
12        printf("I'm the parent.\n");
13        printf("My PID is %d\n", getpid());
14    }
15
16    printf("program terminated.\n");
17 }
```



```
$ ./fork_example_2
before fork ...
result = 1235
I'm the parent.
My PID is 1234
program terminated.
```

PID 1234
(stop)

PID 1235
(waiting)

Creating Processes with fork() System Call

```
1 int main(void) {  
2     int result;  
3     printf("before fork ...\n");  
4     result = fork();  
5     printf("result = %d.\n", result);  
6  
7     if(result == 0) {  
8         printf("I'm the child.\n");  
9         printf("My PID is %d\n", getpid());  
10    }  
11    else {  
12        printf("I'm the parent.\n");  
13        printf("My PID is %d\n", getpid());  
14    }  
15  
16    printf("program terminated.\n");  
17 }
```

```
$ ./fork_example_2  
before fork ...  
result = 1235  
I'm the parent.  
My PID is 1234  
program terminated.  
result = 0
```

Important

For child, the return value of `fork()` is 0.

PID 1234
(stop)

PID 1235
(running)

Creating Processes with fork() System Call

```
1 int main(void) {  
2     int result;  
3     printf("before fork ...\n");  
4     result = fork();  
5     printf("result = %d.\n", result);  
6  
7     if(result == 0) {  
8         printf("I'm the child.\n");  
9         printf("My PID is %d\n", getpid());  
10    }  
11    else {  
12        printf("I'm the parent.\n");  
13        printf("My PID is %d\n", getpid());  
14    }  
15  
16    printf("program terminated.\n");  
17 }
```

```
$ ./fork_example_2  
before fork ...  
result = 1235  
I'm the parent.  
My PID is 1234  
program terminated.  
result = 0  
I'm the child.  
My PID is 1235
```

PID 1234
(stop)

PID 1235
(running)

Creating Processes with fork() System Call

```
1 int main(void) {  
2     int result;  
3     printf("before fork ...\n");  
4     result = fork();  
5     printf("result = %d.\n", result);  
6  
7     if(result == 0) {  
8         printf("I'm the child.\n");  
9         printf("My PID is %d\n", getpid());  
10    }  
11    else {  
12        printf("I'm the parent.\n");  
13        printf("My PID is %d\n", getpid());  
14    }  
15  
16    printf("program terminated.\n");  
17 }
```

```
$ ./fork_example_2  
before fork ...  
result = 1235  
I'm the parent.  
My PID is 1234  
program terminated.  
result = 0  
I'm the child.  
My PID is 1235  
program terminated.  
$ _
```

PID 1234
(stop)

PID 1235
(stop)

fork() System Call

- `fork()` behaves like “cell division”.
 - It creates the child process by **cloning** from the parent process, including all user-space data, e.g.,

Cloned items	Descriptions
Program counter [CPU register]	That’s why they both execute from the same line of code after <code>fork()</code> returns.
Program code [File & Memory]	They are sharing the same piece of code.
Memory	Including local variables, global variables, and dynamically allocated memory.
Opened files [Kernel’s internal]	If the parent has opened a file “fd”, then the child will also have file “fd” opened automatically.

fork() System Call

- fork() does not clone the following...

Distinct items	Parent	Child
Return value of fork()	PID of the child process.	0
PID	Unchanged.	Different, not necessarily be "Parent PID + 1"
Parent process	Unchanged.	Parent.
Running time	Cumulated.	Just created, so should be 0.
[Advanced] File locks	Unchanged.	None.

fork() System Call

- If a process can only duplicate itself and always runs the same program, it's not quite meaningful
 - how can we execute other programs?
- **exec()**
 - The **exec*()** system call family.

exec()

- `execl()` – a member of the `exec` system call family (`execl`, `execle`, `execlp`, `execv`, `execve`, `execvp`).

```
int main(void) {  
    printf("before execl ...\n");  
    execl("/bin/ls", "/bin/ls", NULL);  
    printf("after execl ...\n");  
    return 0;  
}
```

```
$ ./exec_example  
before execl ...
```

Arguments of the `execl()` call

1st argument: the program name, `"/bin/ls"` in the example.

2nd argument: `argument[0]` to the program.

3rd argument: `argument[1]` to the program.

exec()

- `execl()` – a member of the `exec` system call family (and the family has 6 members).

```
int main(void) {  
    printf("before execl ...\n");  
    → execl("/bin/ls", "/bin/ls", NULL);  
    printf("after execl ...\n");  
    return 0;  
}
```

```
$/exec_example  
before execl ...  
exec_example  
exec_example.c
```

What is the output?

The same as the output of running “ls” in the shell.

exec()

- Example #1: run the command **"/bin/ls"**

```
execl("/bin/ls", "/bin/ls", NULL);
```

Argument Order	Value in above example	Description
1	"/bin/ls"	The file that the programmer wants to execute.
2	"/bin/ls"	When the process switches to "/bin/ls" , this string is the program argument[0] .
3	NULL	This states the end of the program argument list.

exec()

- Example #2: run the command **"/bin/ls -l"**

```
execl("/bin/ls", "/bin/ls", "-l", NULL);
```

Argument Order	Value in above example	Description
1	"/bin/ls"	The file that the programmer wants to execute.
2	"/bin/ls"	When the process switches to "/bin/ls" , this string is the program argument[0] .
3	"-l"	When the process switches to "/bin/ls" , this string is the program argument[1] .
4	NULL	This states the end of the program argument list.

exec()

- The exec system call family is not simply a function that “invokes” a command.

```
int main(void) {  
    printf("before execl ...\n");  
    execl("/bin/ls", "/bin/ls", NULL);  
    printf("after execl ...\n");  
    return 0;  
}
```

WHAT?!
The shell prompt appears!

```
./exec_example  
before execl ...  
exec_example  
exec_example.c  
$ _
```


The output says:

- (1) The gray code block **is not reached!**
- (2) The process is **terminated!**

WHY IS THAT?!

exec()

- The exec system call family is not simply a function that “invokes” a command.

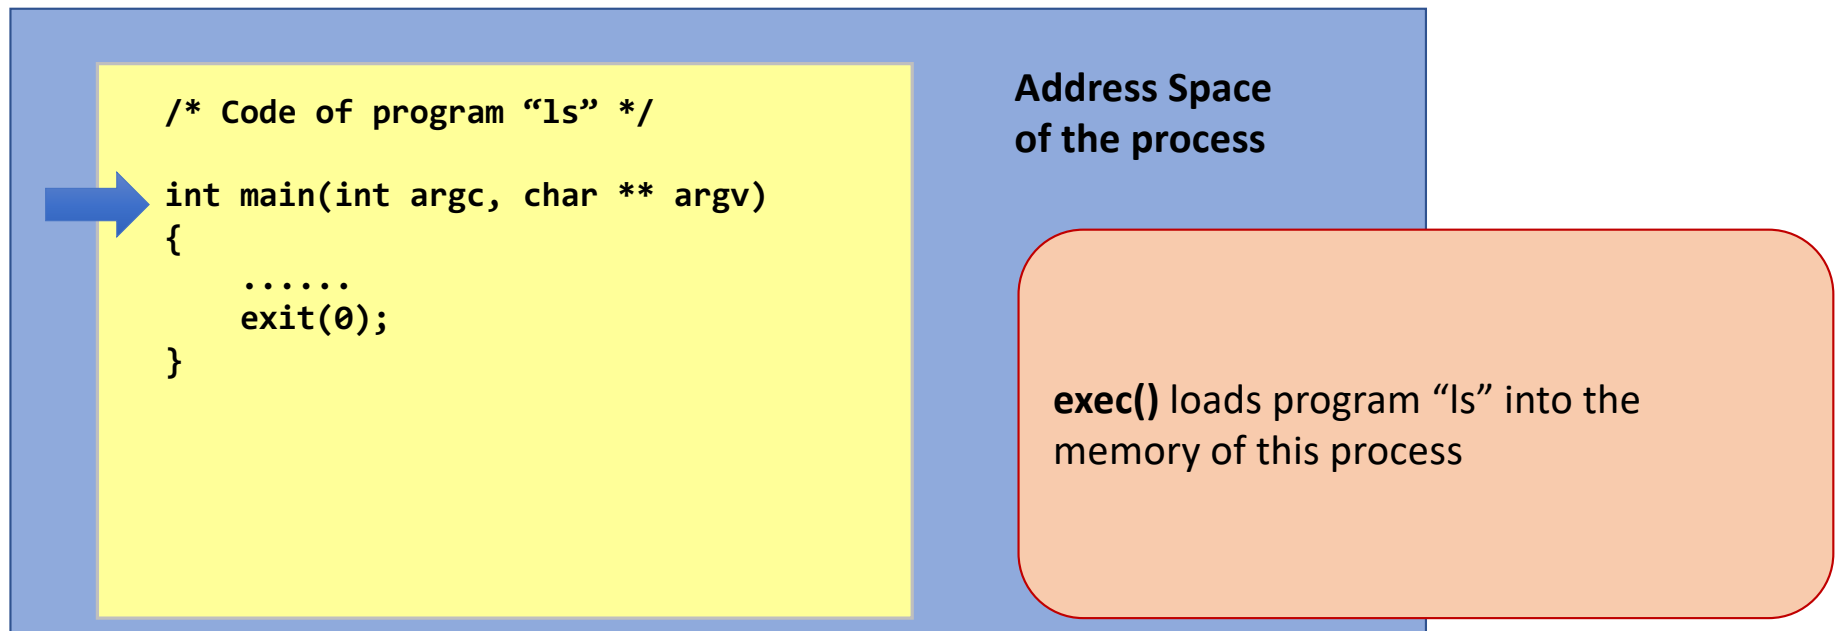


```
/* code of program exec_example */  
int main(void) {  
    printf("before execl ...\n");  
    execl("/bin/ls", "/bin/ls", NULL);  
    printf("after execl ...\n");  
    return 0;  
}
```

**Address Space
of the process**

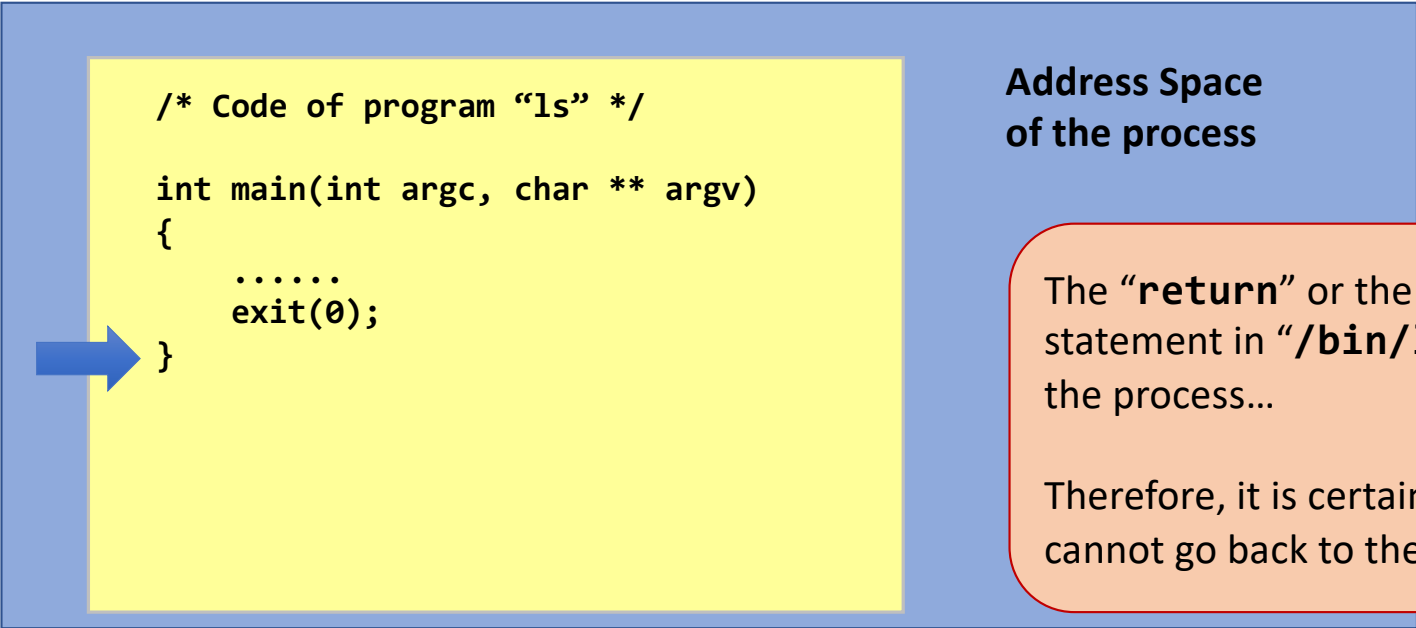
exec()

- The `exec` system call family is not simply a function that “invokes” a command.



exec()

- The `exec` system call family is not simply a function that “invokes” a command.



The diagram illustrates the address space of a process. It consists of a large blue rectangle representing the entire address space. Inside this rectangle, on the left, is a yellow rectangle representing the code of the program 'ls'. A blue arrow points from the left edge of the yellow rectangle to its right edge, indicating the flow of execution. To the right of the yellow rectangle, within the blue area, is a red-bordered rounded rectangle containing text about process termination. The yellow rectangle contains the following C code:

```
/* Code of program "ls" */  
  
int main(int argc, char ** argv)  
{  
    .....  
    exit(0);  
}
```

Address Space
of the process

The “**return**” or the “**exit()**” statement in “**/bin/ls**” will terminate the process...

Therefore, it is certain that the process cannot go back to the old program!

exec() Summary

- The process is changing the code that is executing and never returns to the original code.
 - The last two lines of codes are therefore not executed.
- The process that calls an exec* system call will replace user-space info, e.g.,
 - Program Code
 - Memory: local variables, global variables, and dynamically allocated memory;
 - Register value: e.g., the program counter;
- But, the kernel-space info of that process is preserved, including:
 - PID;
 - Process relationship;
 - etc.

CPU Scheduler and fork()

```
1 int main(void) {
2     int result;
3     printf("before fork ...\n");
4     result = fork();
5     printf("result = %d.\n", result);
6
7     if(result == 0) {
8         printf("I'm the child.\n");
9         printf("My PID is %d\n", getpid());
10    }
11    else {
12        printf("I'm the parent.\n");
13        printf("My PID is %d\n", getpid());
14    }
15
16    printf("program terminated.\n");
17 }
```

Parent return
from fork() first

```
$ ./fork_example_2
before fork ...
result = 1235
I'm the parent.
My PID is 1234
program terminated.
result = 0
I'm the child.
My PID is 1235
program terminated.
$ _
```

Child return
from fork() first

```
$ ./fork_example_2
before fork ...
result = 0
I'm the child.
My PID is 1235
result = 1235
program terminated.
I'm the parent.
My PID is 1234
program terminated.
$ _
```

wait(): Sync Parent with Child

```
1 int main(void) {
2     int result;
3     printf("before fork ...\n");
4     result = fork();
5     printf("result = %d.\n", result);
6
7     if(result == 0) {
8         printf("I'm the child.\n");
9         printf("My PID is %d\n", getpid());
10    }
11    else {
12        printf("I'm the parent.\n");
13        wait(NULL);
14        printf("My PID is %d\n", getpid());
15    }
16
17    printf("program terminated.\n");
18 }
```

Parent return
from fork() first


```
$ ./fork_example_2
before fork ...
result = 1235
I'm the parent.
result = 0
I'm the child.
My PID is 1235
program terminated.
My PID is 1234
program terminated.
$ _
```

Child return
from fork() first

```
$ ./fork_example_2
before fork ...
result = 0
I'm the child.
My PID is 1235
result = 1235
program terminated.
I'm the parent.
My PID is 1234
program terminated.
$ _
```

wait()

- wait() suspends the calling process to **waiting**
- wait() returns when
 - one of its child processes changes from running to terminated.
- Return immediately (i.e., does nothing) if
 - It has no children
 - Or a child terminates before the parent calls wait for

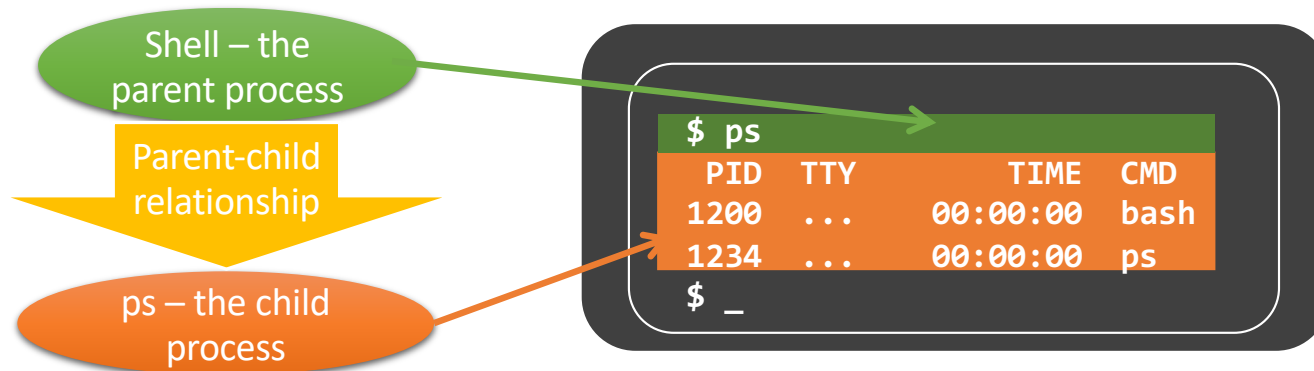


wait() v.s. waitpid()

- **wait()**
 - Wait for any one of the child processes
 - Detect child termination only
- **waitpid()**
 - Depending on the parameters, waitpid() will wait for a particular child only
 - Depending on the parameters, waitpid() can detect different status changes of the child (resume/stop by a signal)

Implement Shell with fork(), exec(), and wait()

- A shell is a CLI
 - Bash in linux
 - invokes a function fork() to create a new process
 - Ask the the child process to exec() the target program
 - Use wait() to wait until the child process terminates

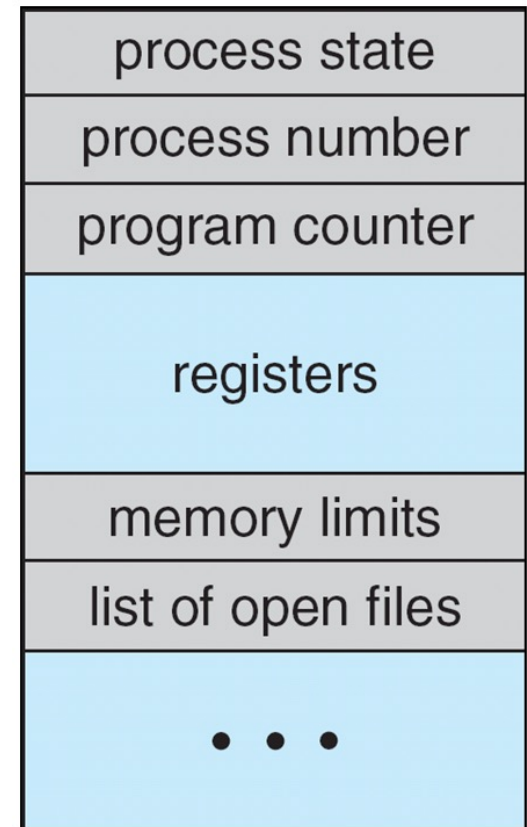


Processes: Kernel View

Process Control Block (PCB)

Information associated with each process

- Process state
- Program counter
- CPU registers
- CPU scheduling information
- Memory-management information
- Accounting information
- I/O status information



PCB Example: uCore

```
/* kern/process/proc.h in ucore */

struct proc_struct {
    enum proc_state state;           // Process state
    int pid;                         // Process ID
    int runs;                        // the running times of Process
    uintptr_t kstack;                // Process kernel stack
    volatile bool need_resched;      // bool value: need to be rescheduled to release CPU?
    struct proc_struct *parent;      // the parent process
    struct mm_struct *mm;            // Process's memory management field
    struct context context;          // Switch here to run process
    struct trapframe *tf;            // Trap frame for current interrupt
    uintptr_t cr3;                   // CR3 register: the base addr of Page Directroy Table(PDT)
    uint32_t flags;                  // Process flag
    char name[PROC_NAME_LEN + 1];   // Process name
    list_entry_t list_link;          // Process link list
}
```

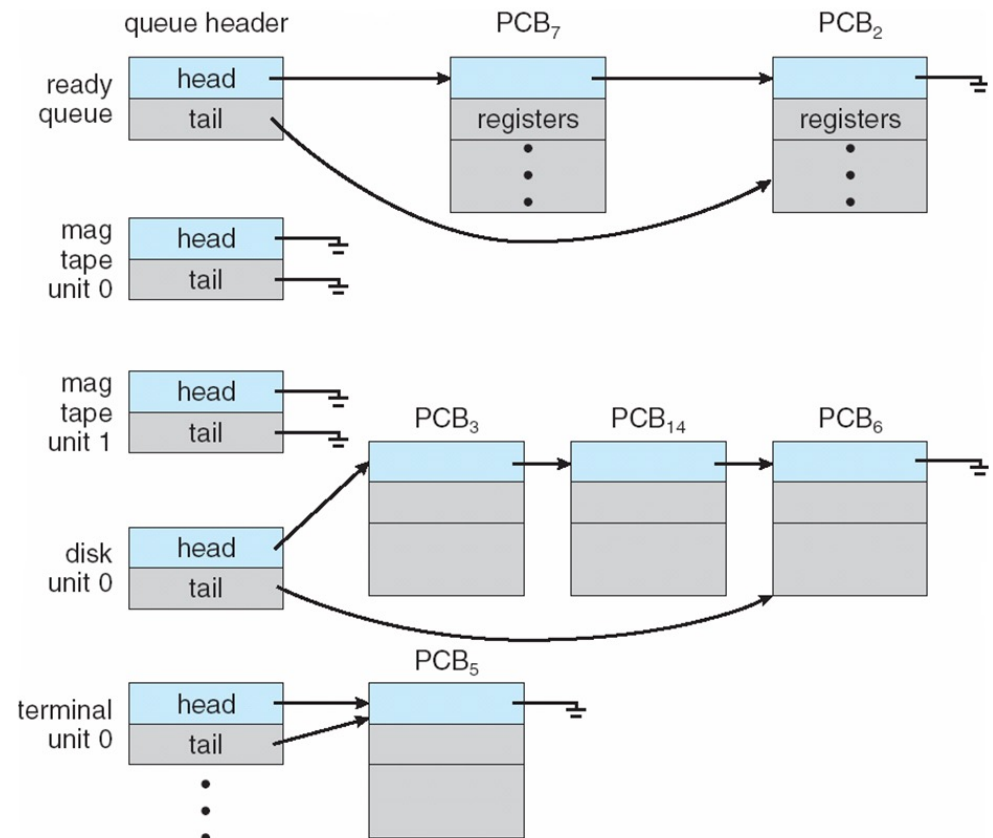
PCB Example: uCore

```
/* kern/process/proc.h in ucore */

list_entry_t hash_link;           // Process hash list
int exit_code;                    // exit code (be sent to parent proc)
uint32_t wait_state;              // waiting state
struct proc_struct *cptr, *yptr, *optr; // relations between processes
struct run_queue *rq;             // running queue contains Process
list_entry_t run_link;            // the entry linked in run queue
int time_slice;                   // time slice for occupying the CPU
struct files_struct *filesp;      // the file related info of process
};
```

Ready Queue And I/O Device Queues

- PCBs are linked in multiple queues
 - Ready queue contains all processes in the ready state (to run on this CPU)
 - Device queue contains processes waiting for I/O events from this device
 - Process may migrate among these queues

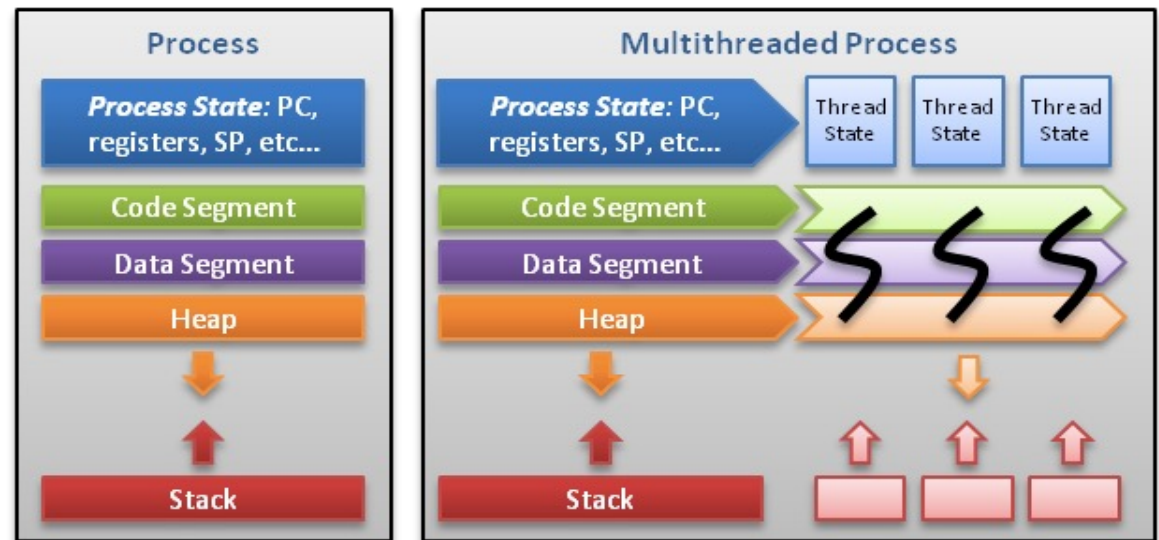


Threads

- One process may have more than one threads
 - A single-threaded process performs a single thread of execution
 - A multi-threaded process performs multiple threads of execution “concurrently”, thus allowing short response time to user’s input even when the main thread is busy
- PCB is extended to include information about each thread

Process and Thread

- Single threaded process and multi-threaded process



Threads contain only necessary information, such as a stack (for local variables, function arguments, return values), a copy of the registers, program counter and any thread-specific data to allow them to be scheduled individually. Other data is shared within the process between all threads.

Switching Between Processes

- Once a process runs on a CPU, it only gives back the control of a CPU
 - when it makes a system call
 - when it raises an exception
 - when an interrupt occurs
- What if none of these would happen for a long time?
 - Cooperative scheduling: OS will have to wait
 - Early Macintosh OS, old Alto system
 - Non-cooperative scheduling: timer interrupts
 - Modern operating systems

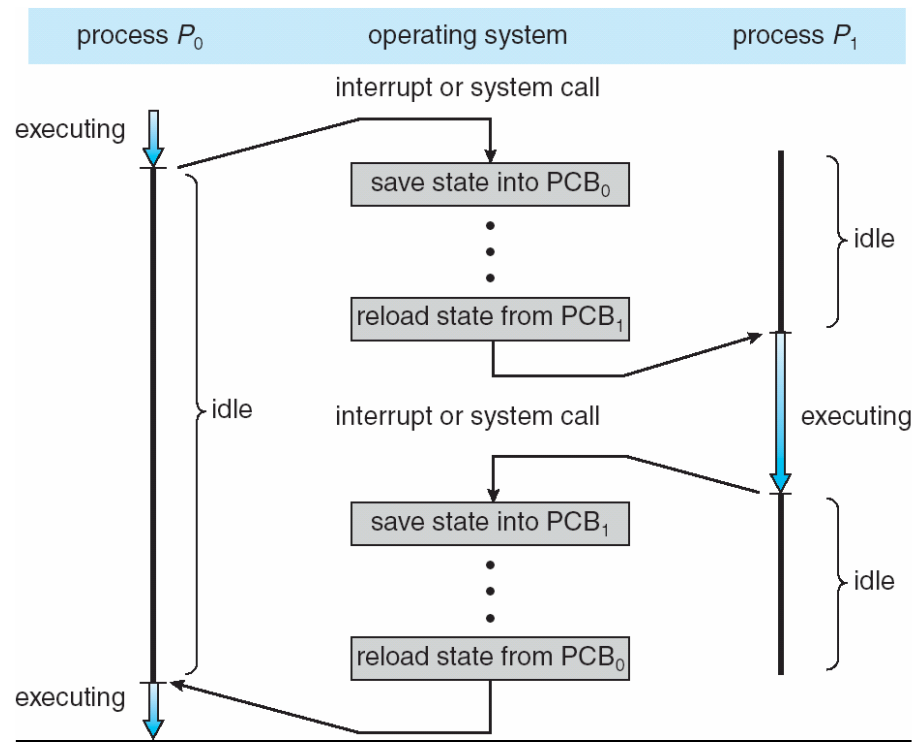
Switching Between Processes (Cont'd)

- When OS kernel regains the control of CPU
 - It first completes the task
 - Serve system call, or
 - Handle interrupt/exception
 - It then decides which process to run next
 - by asking its **CPU scheduler**
 - How does it make decisions?
 - More about CPU scheduler later
 - It performs a **context switch** if the soon-to-be-executing process is different from the previous one

Context Switch

- During context switch, the system must save the state of the old process and load the saved state for the new process
- Context of a process is represented in the PCB
- The time used to do context switch is an overhead of the system; the system does no useful work while switching
 - Time of context switch depends on hardware support
 - Context switch cannot be too frequent

Context Switch (Cont'd)



Context Switch: uCore

```
/* kern/schedule/sched.c */
void schedule(void) {
    bool intr_flag;
    struct proc_struct *next;
    local_intr_save(intr_flag);
    {
        if (current->state == PROC_RUNNABLE)
            sched_class_enqueue(current);

        if ((next = sched_class_pick_next()) != NULL)
            sched_class_dequeue(next);

        if (next != current)
            proc_run(next);
    }
    local_intr_restore(intr_flag);
}
```

```
/* kern/process/proc.c */

void proc_run(struct proc_struct *proc) {
    if (proc != current) {
        bool intr_flag;
        struct proc_struct *prev = current, *next = proc;
        local_intr_save(intr_flag);
        {
            current = proc;
            lcr3(next->cr3);
            switch_to(&(prev->context), &(next->context));
        }
        local_intr_restore(intr_flag);
    }
}
```

Context Switch: uCore (Cont'd)

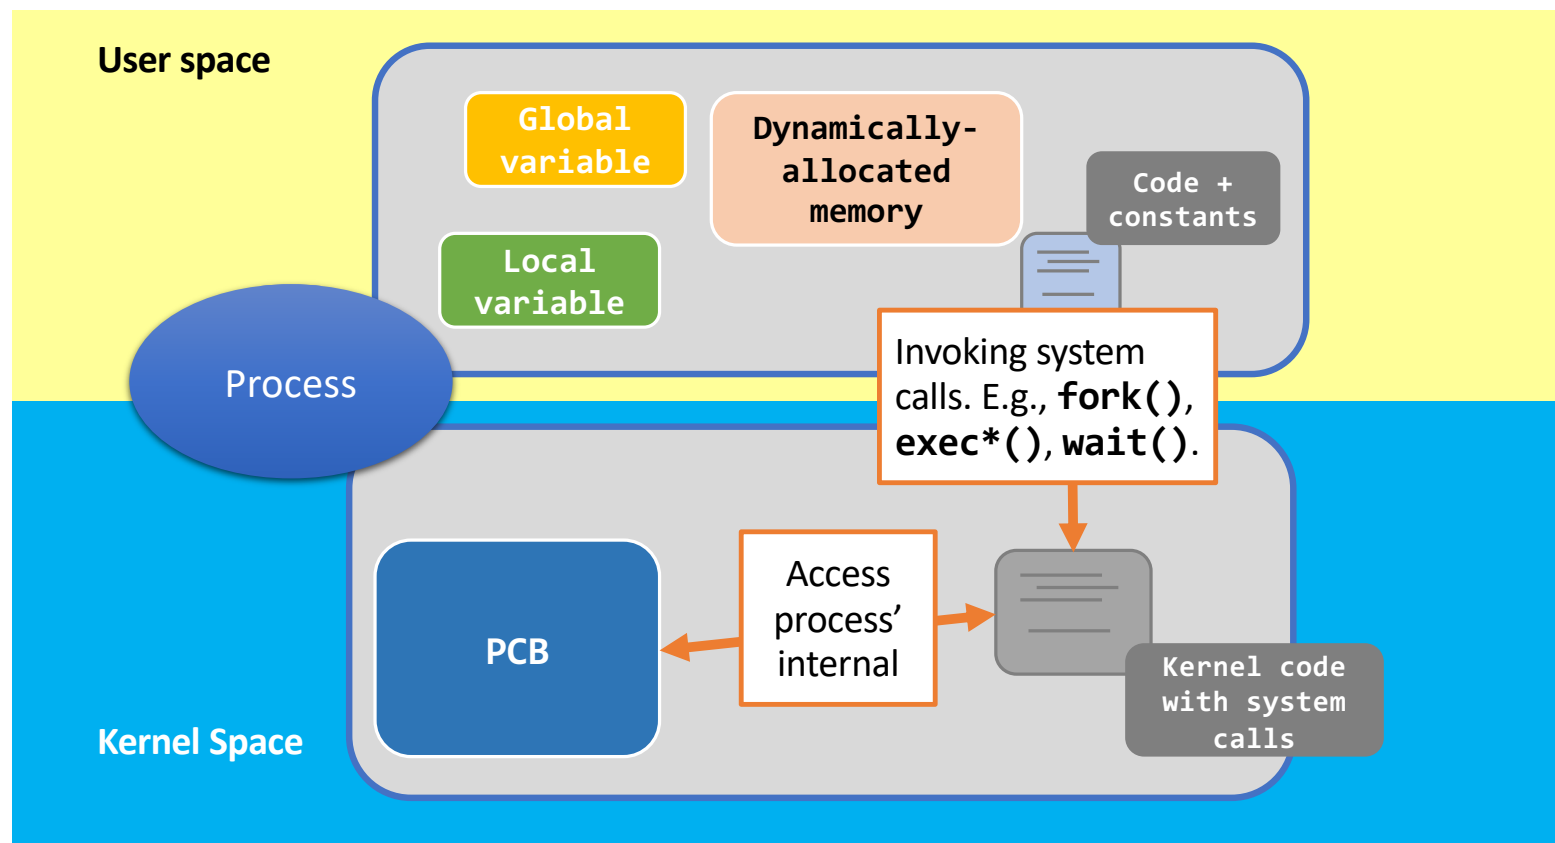
```
/* kern/process/switch.S */
.globl switch_to
switch_to:
    # save from's registers
    STORE ra, 0*REGBYTES(a0)
    STORE sp, 1*REGBYTES(a0)
    STORE s0, 2*REGBYTES(a0)
    STORE s1, 3*REGBYTES(a0)
    STORE s2, 4*REGBYTES(a0)
    STORE s3, 5*REGBYTES(a0)
    STORE s4, 6*REGBYTES(a0)
    STORE s5, 7*REGBYTES(a0)
    STORE s6, 8*REGBYTES(a0)
    STORE s7, 9*REGBYTES(a0)
    STORE s8, 10*REGBYTES(a0)
    STORE s9, 11*REGBYTES(a0)
    STORE s10, 12*REGBYTES(a0)
    STORE s11, 13*REGBYTES(a0)
```

```
# restore to's registers
LOAD ra, 0*REGBYTES(a1)
LOAD sp, 1*REGBYTES(a1)
LOAD s0, 2*REGBYTES(a1)
LOAD s1, 3*REGBYTES(a1)
LOAD s2, 4*REGBYTES(a1)
LOAD s3, 5*REGBYTES(a1)
LOAD s4, 6*REGBYTES(a1)
LOAD s5, 7*REGBYTES(a1)
LOAD s6, 8*REGBYTES(a1)
LOAD s7, 9*REGBYTES(a1)
LOAD s8, 10*REGBYTES(a1)
LOAD s9, 11*REGBYTES(a1)
LOAD s10, 12*REGBYTES(a1)
LOAD s11, 13*REGBYTES(a1)
```


```
ret
```

fork(), exec(), wait()
Kernel View


Recall: fork(), exec(), and wait()



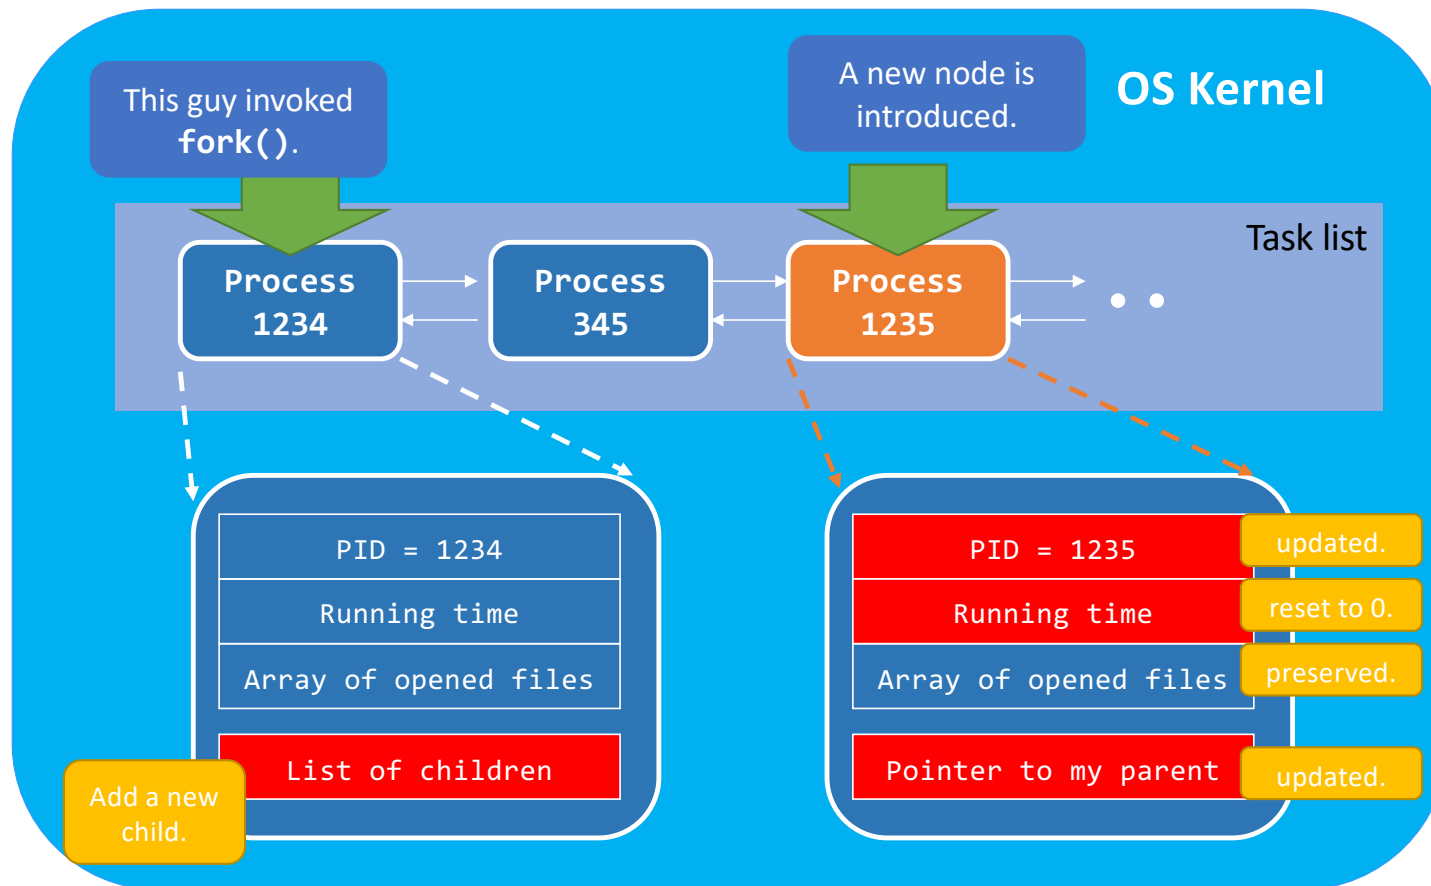
Fork() in User Mode




fork(): Kernel View



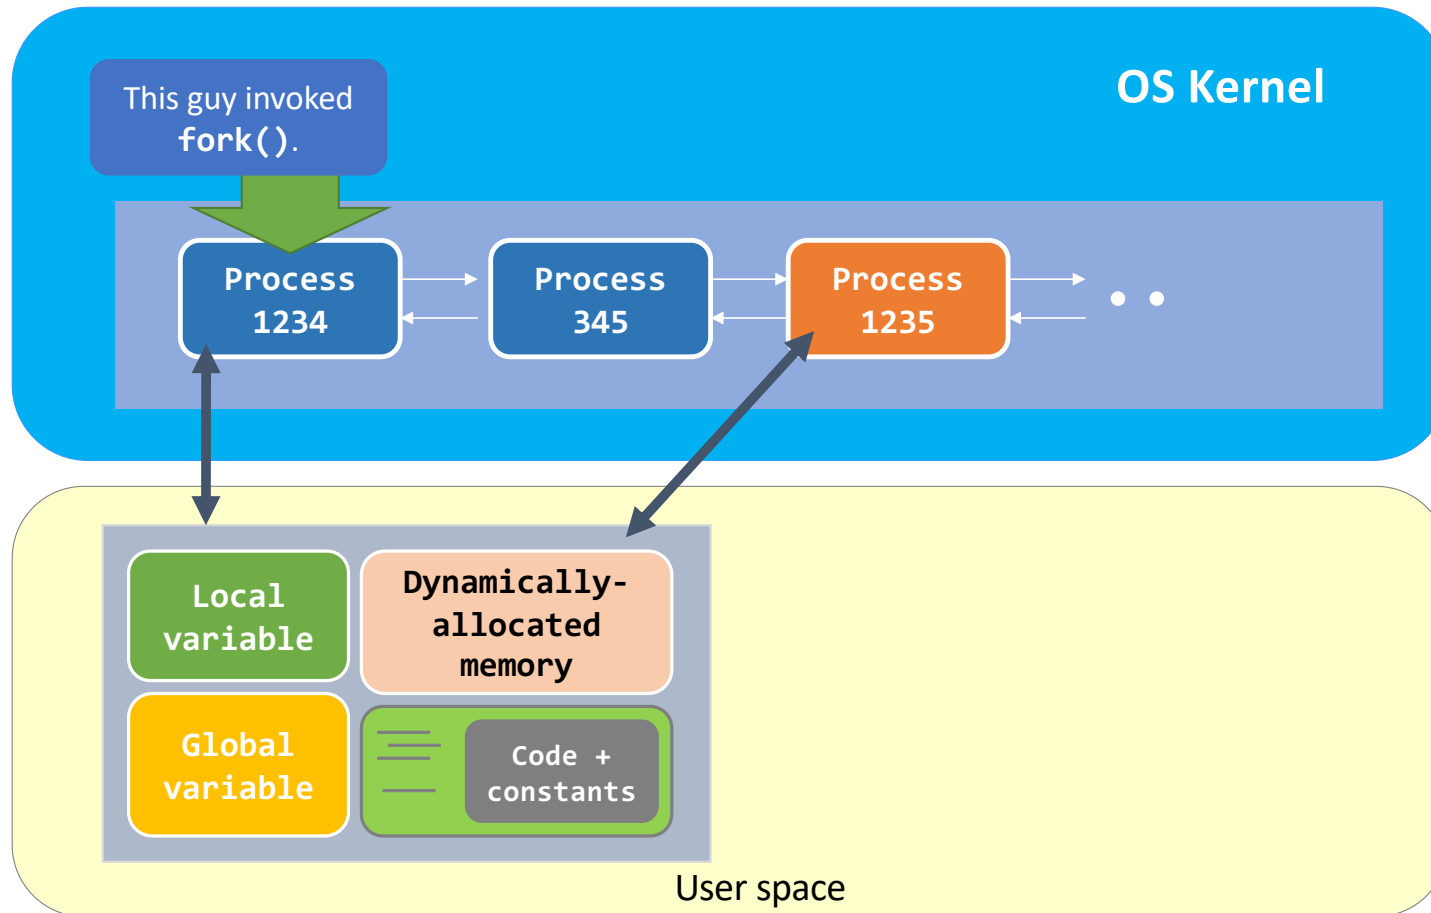
fork(): Kernel View




Case 1: Duplicate Address Space



Case 2: Copy on Write



fork(): Kernel View



fork(): Opened Files


- Array of opened files contains:

Array Index	Description
0	Standard Input Stream; FILE *stdin;
1	Standard Output Stream; FILE *stdout;
2	Standard Error Stream; FILE *stderr;
3 or beyond	Storing the files you opened, e.g., fopen() , open() , etc.

- That's why a parent process shares the same terminal output stream as the child process.


fork(): Opened Files

- What if two processes, sharing the same opened file, write to that file together?

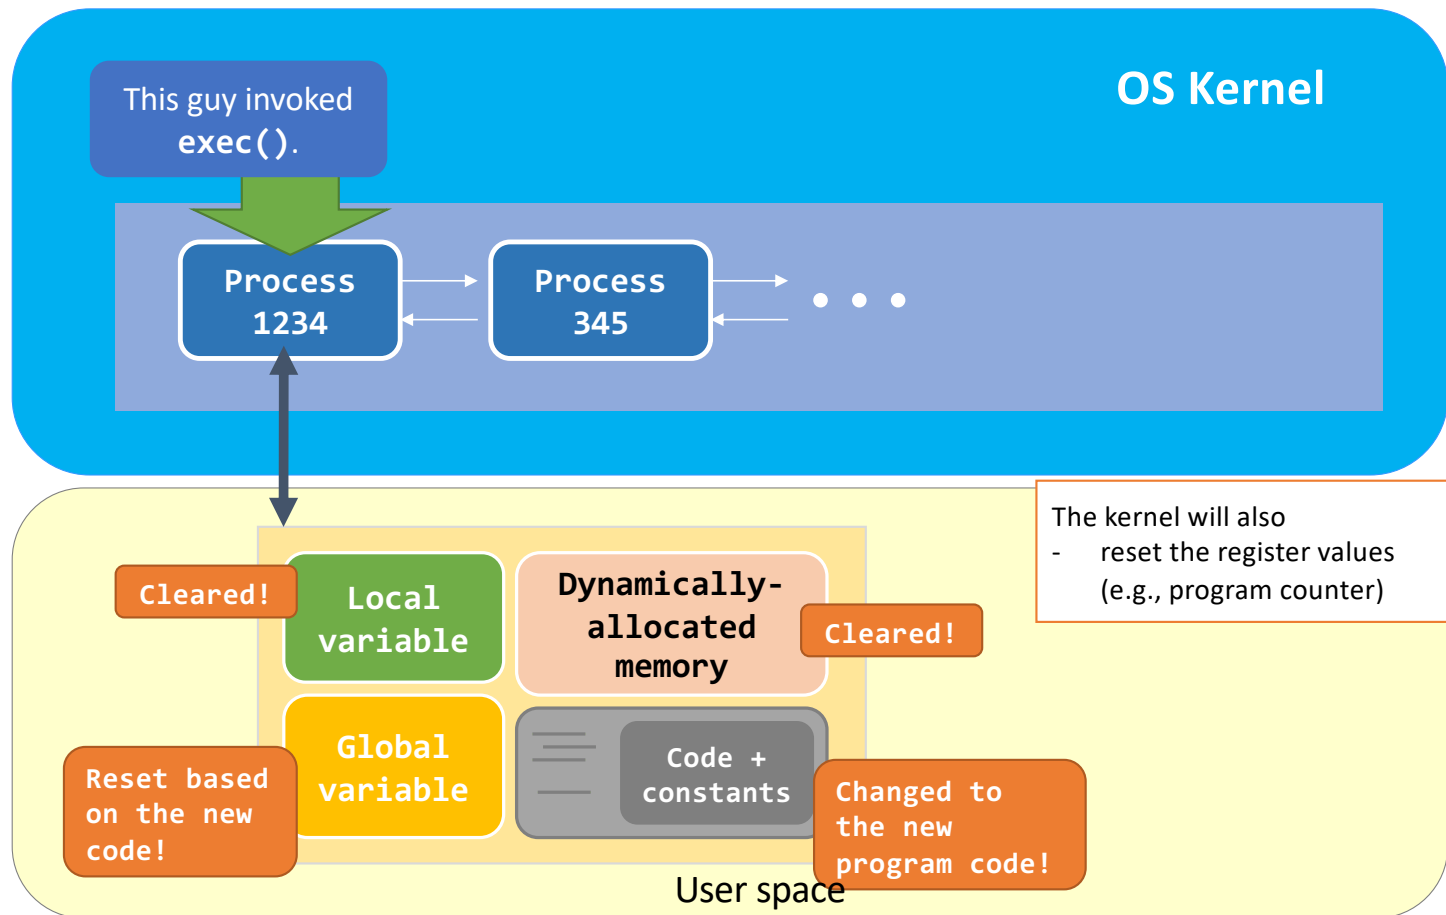


Let's see what will happen when the program finishes running!

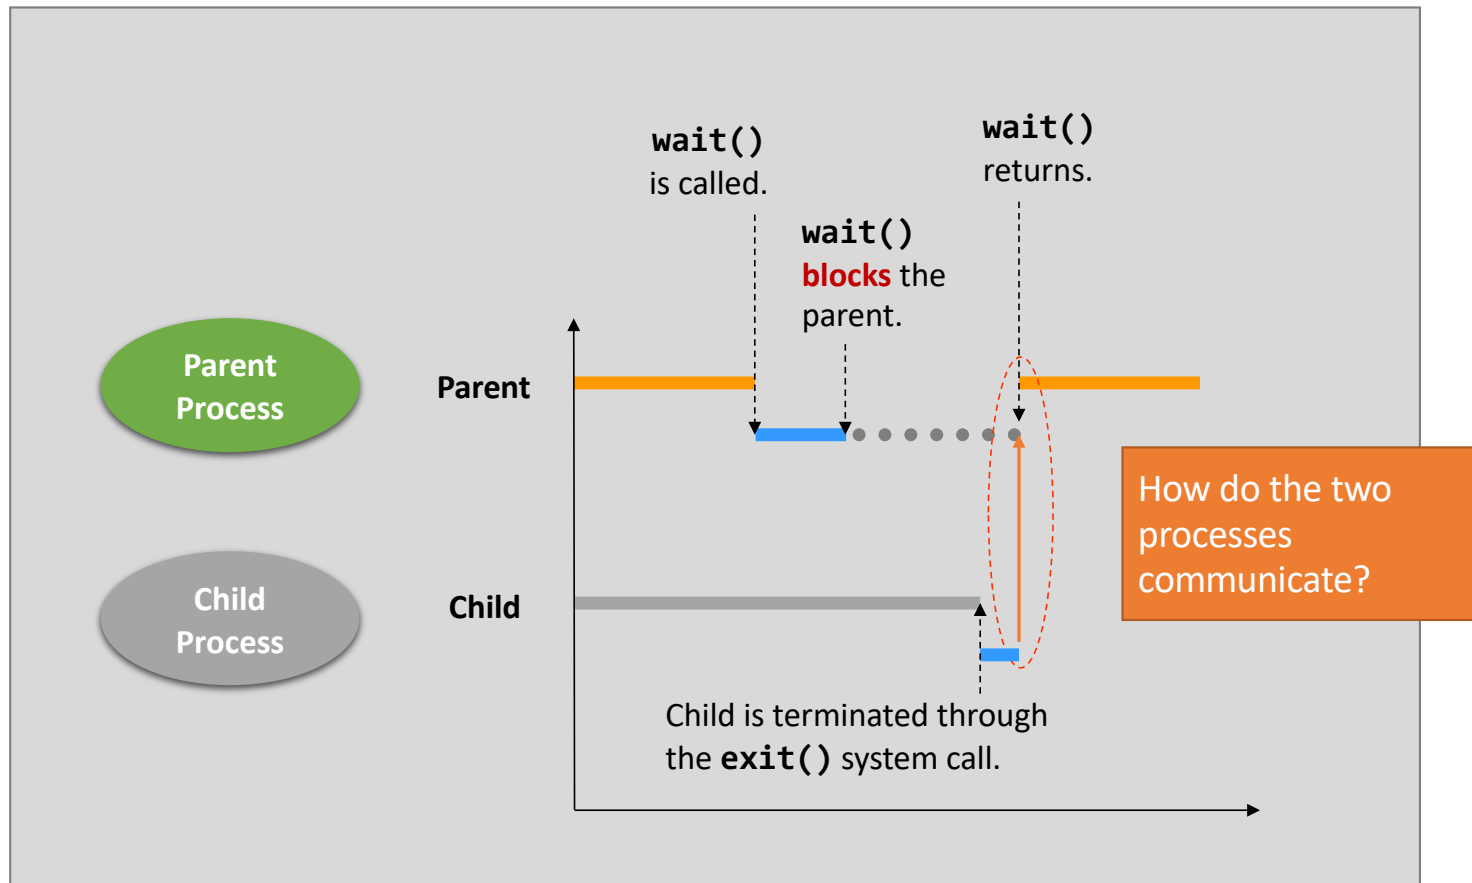
exec() in the User Mode



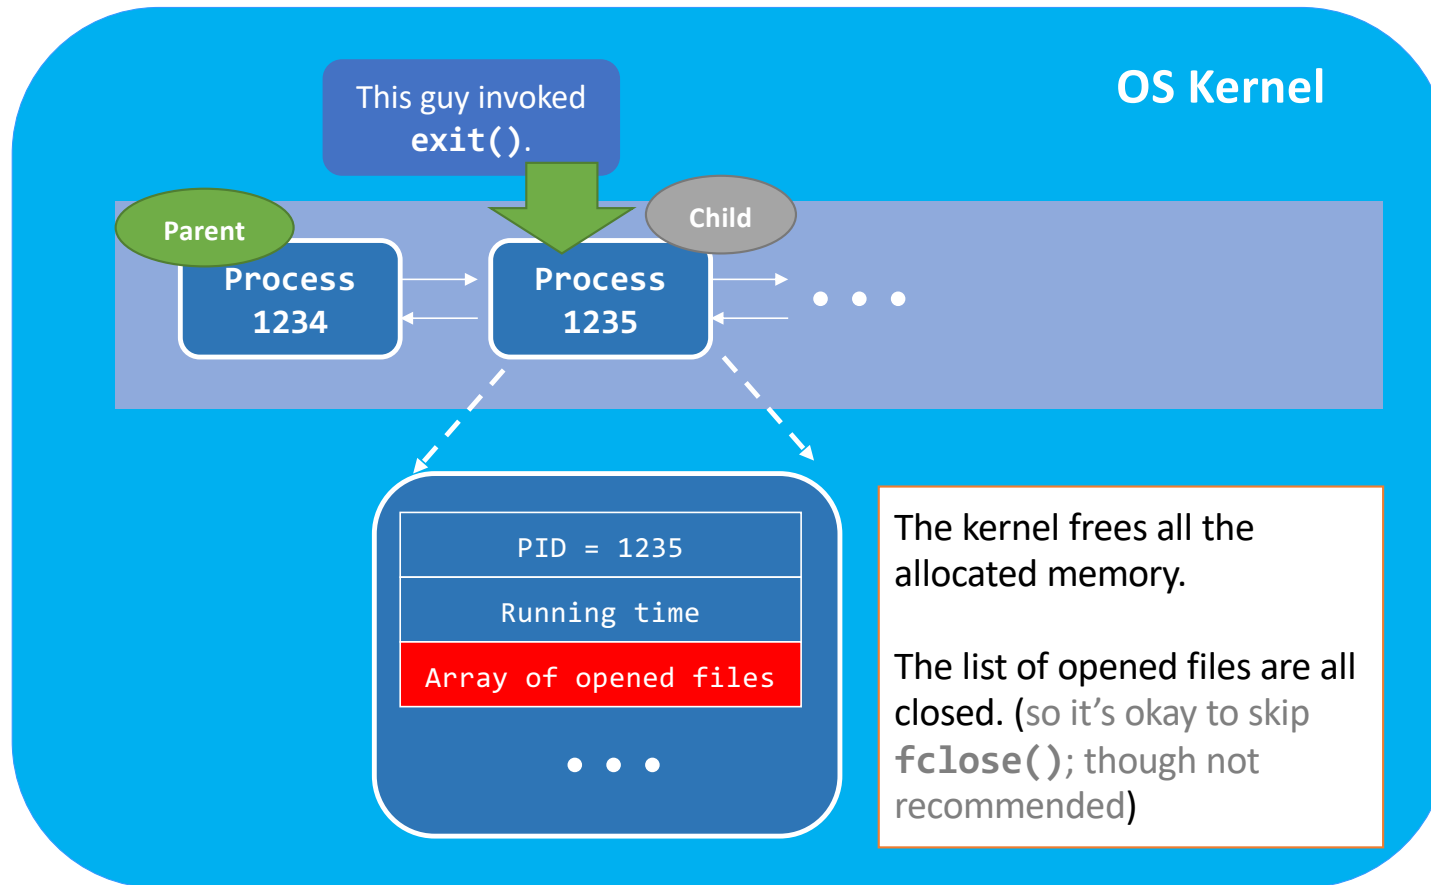
exec(): Kernel View



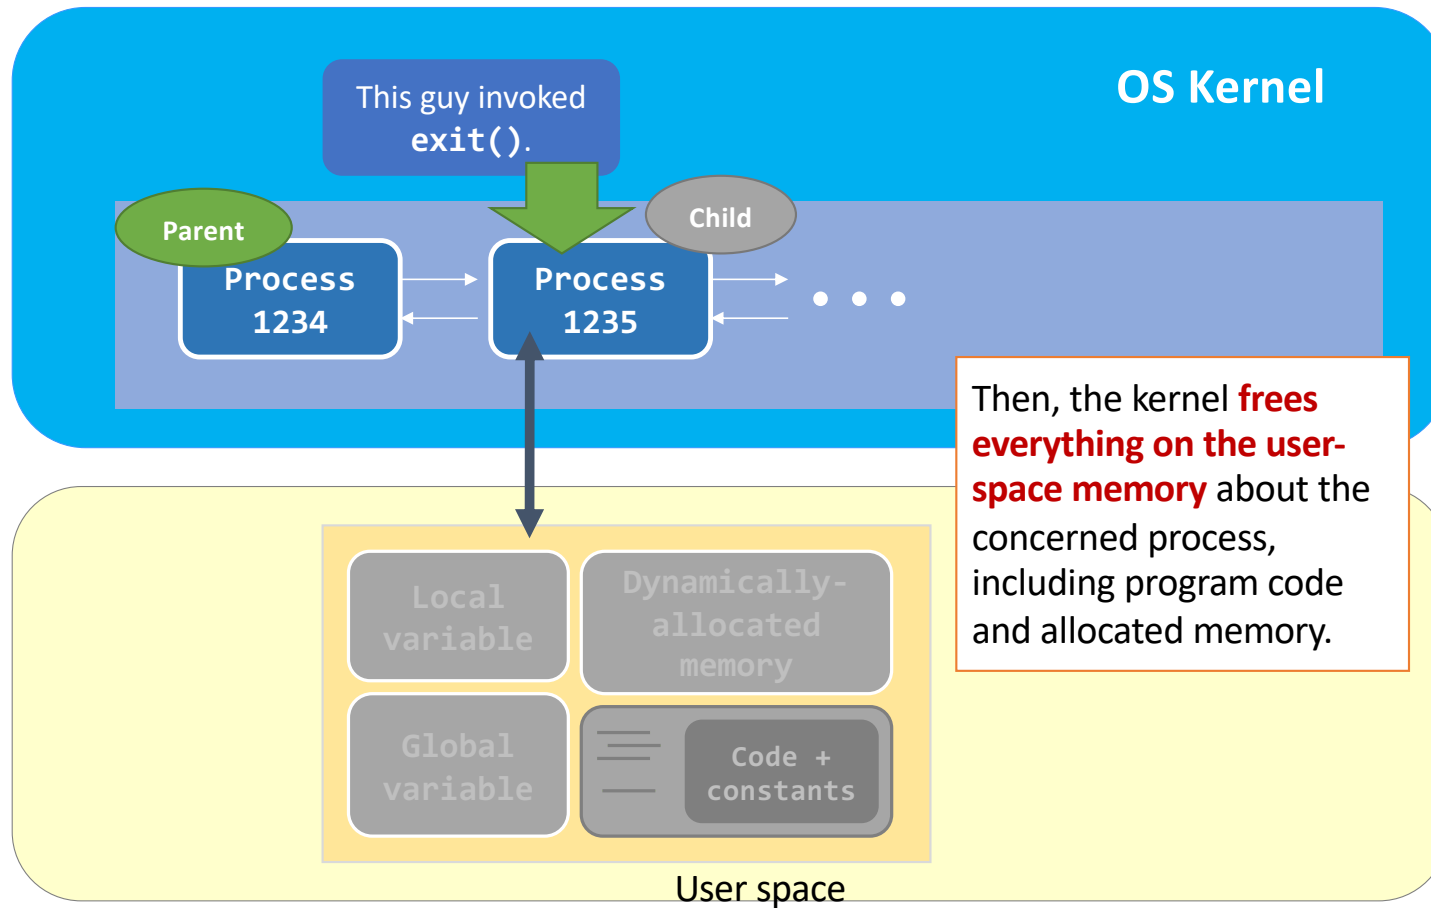
wait() and exit()



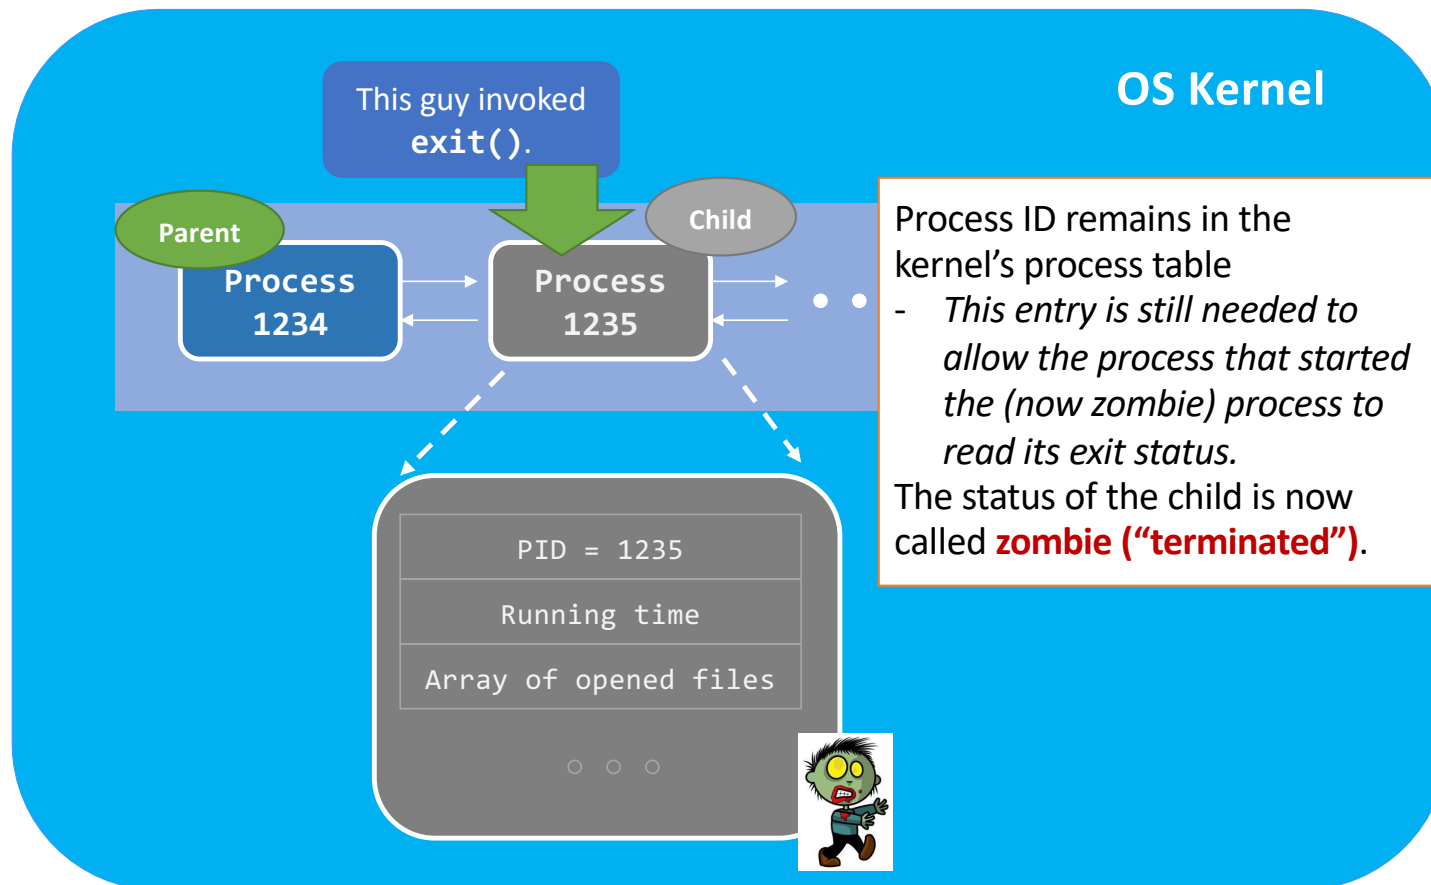
exit(): Kernel View



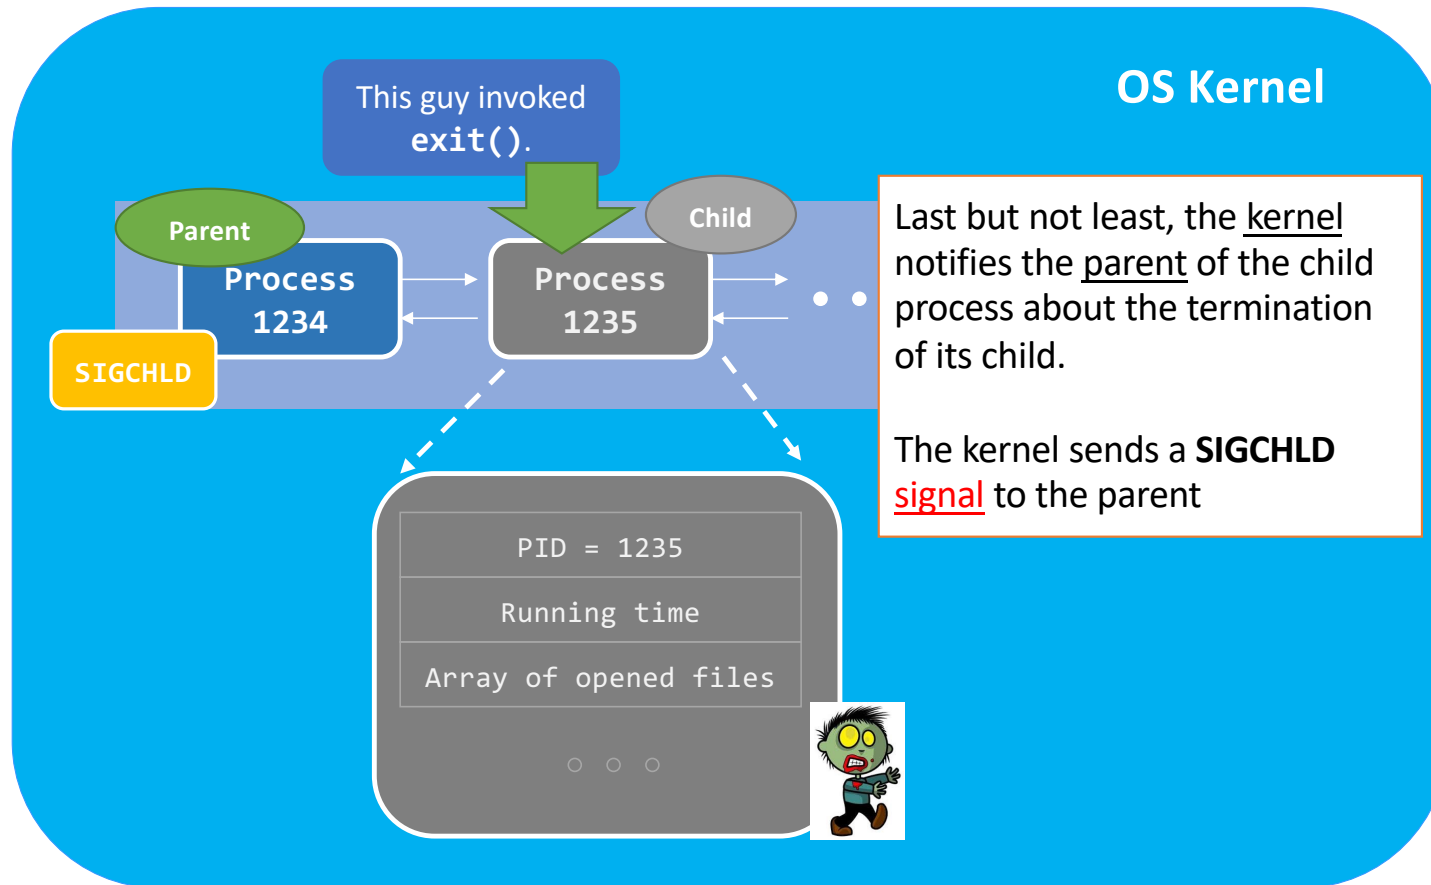
exit(): Kernel View



exit(): Kernel View



exit(): Kernel View




exit(): Summary

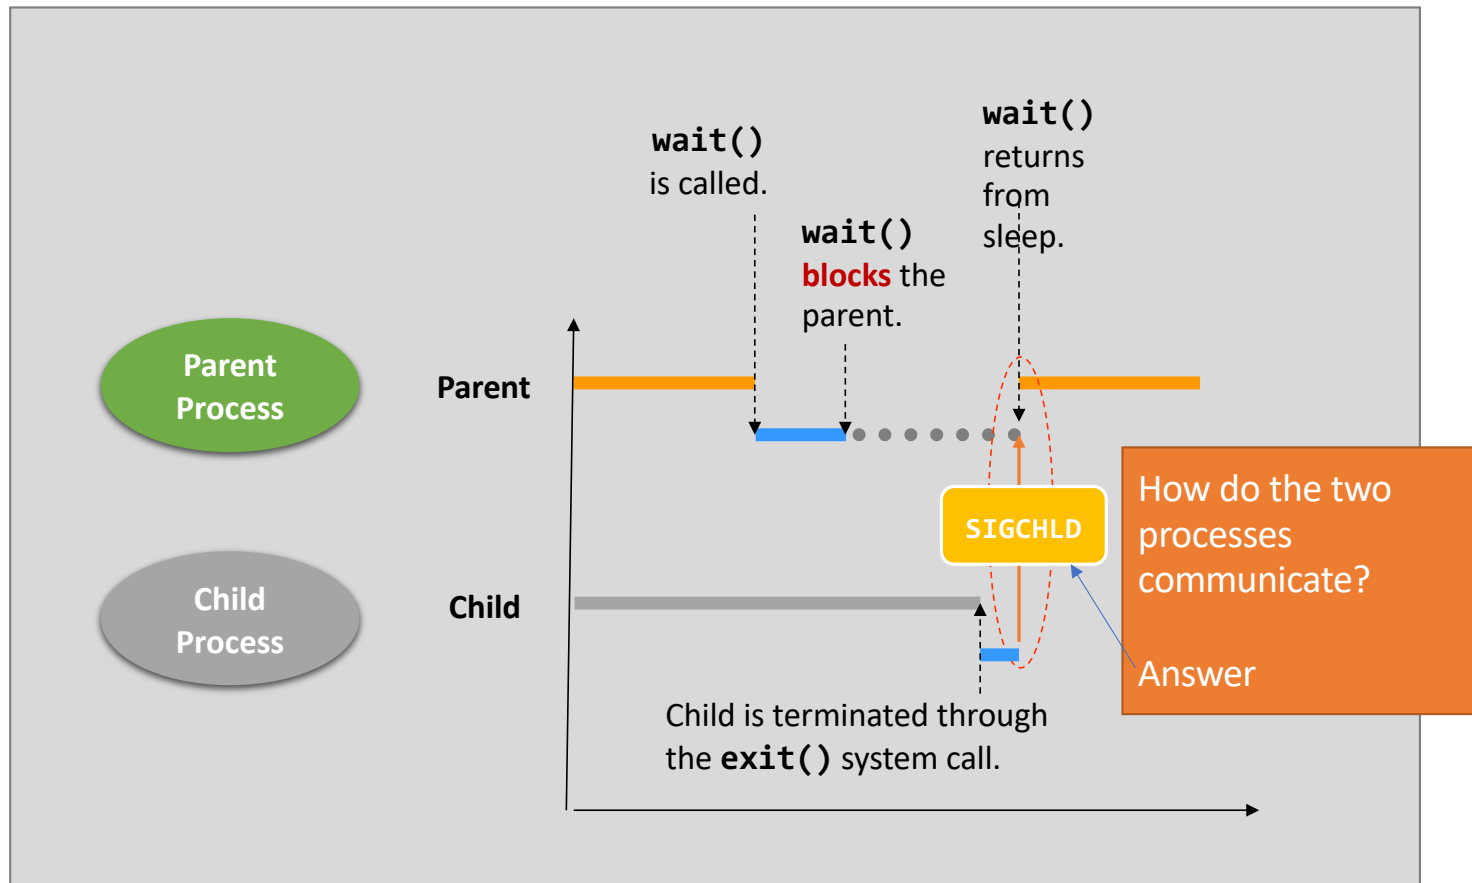
Step (1) Clean up most of the allocated kernel-space memory (e.g., process's running time info).

Step (2) Clean up the exit process's user-space memory.

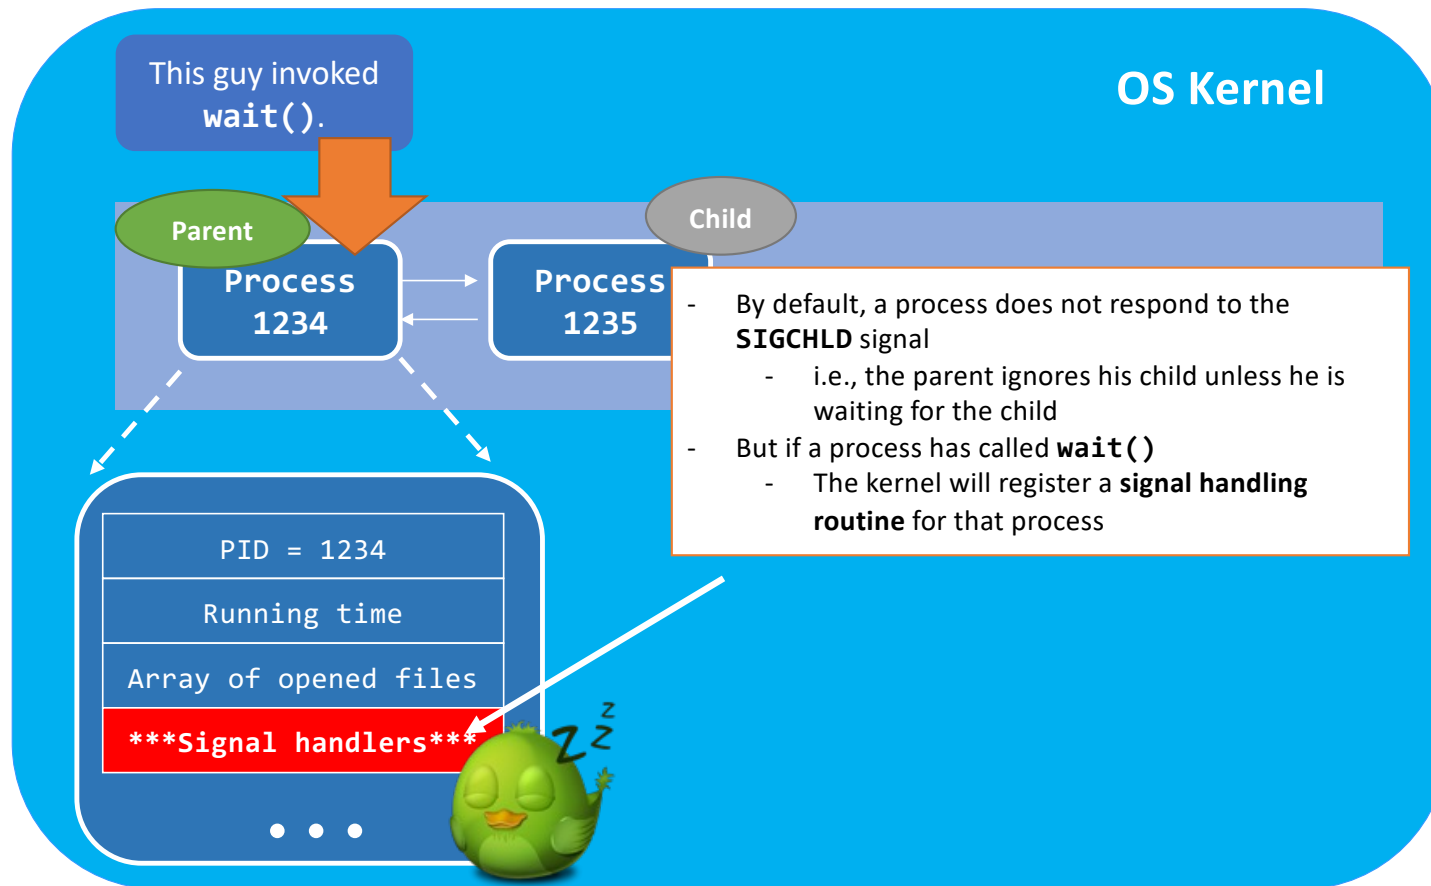
Step (3) Notify the parent with SIGCHLD.



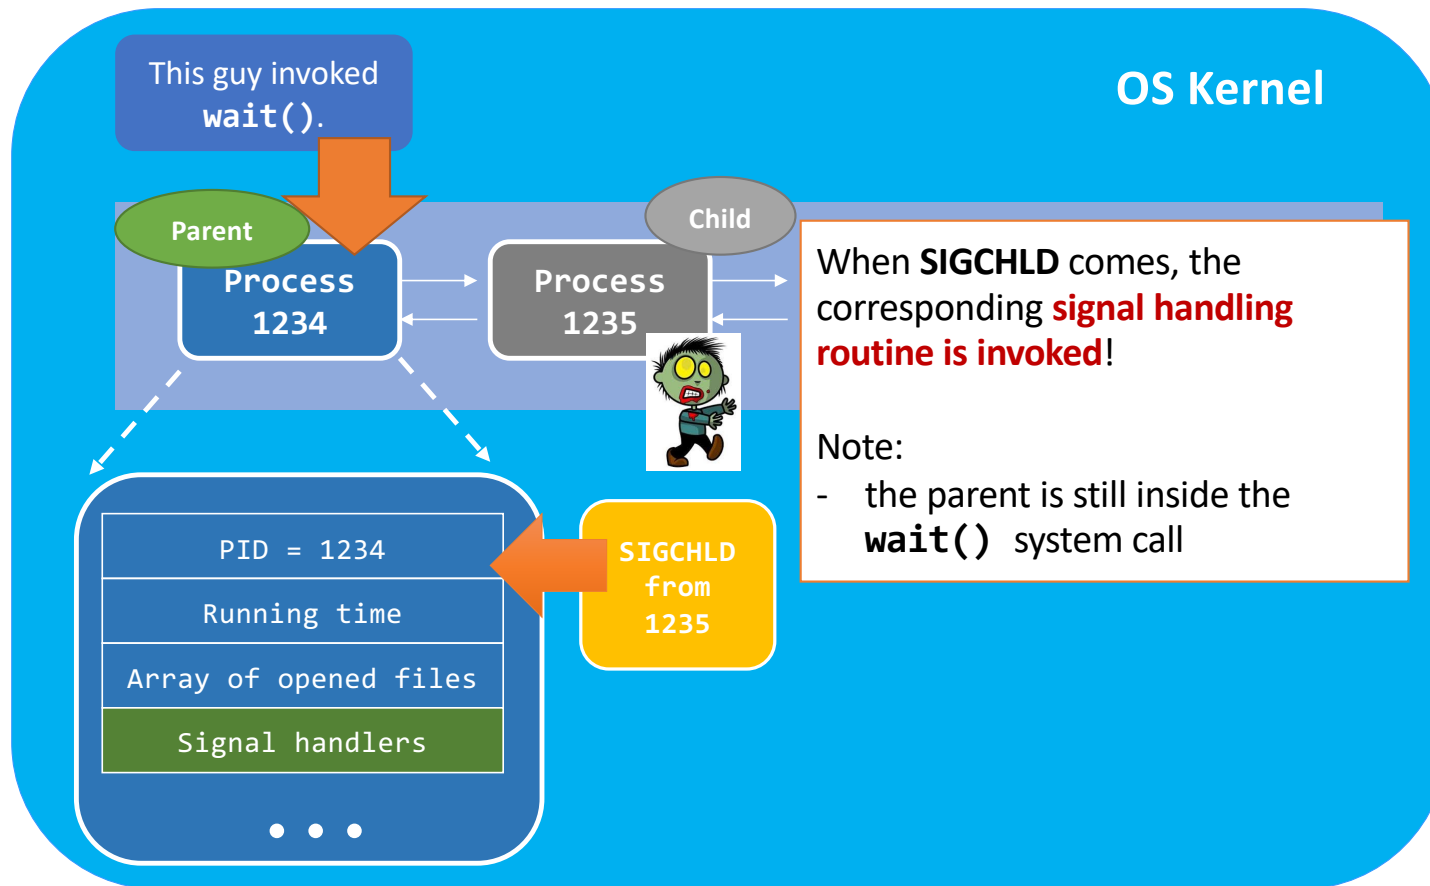
wait() and exit()



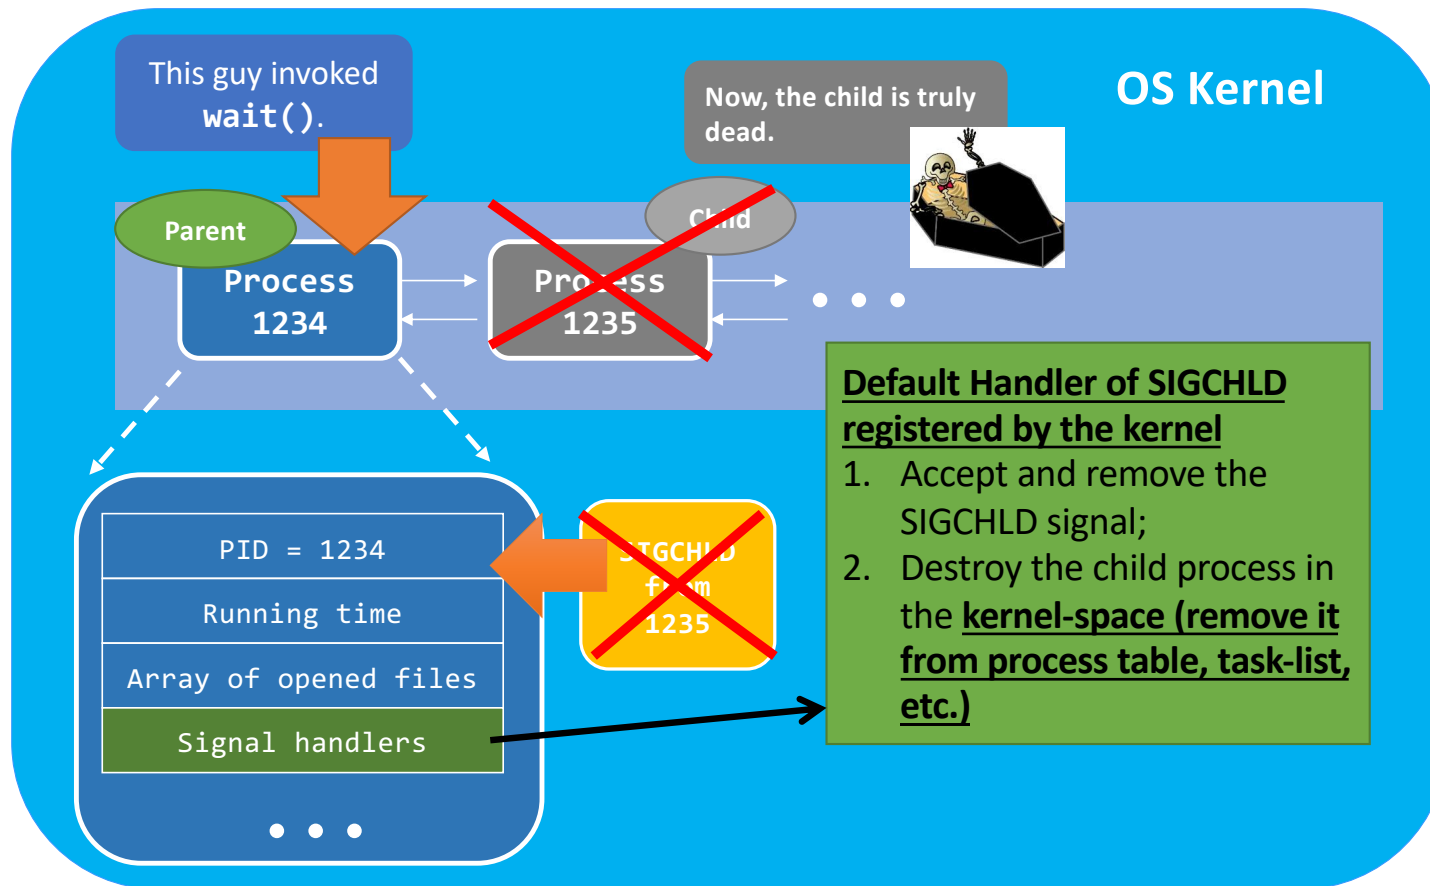
wait() Kernel View



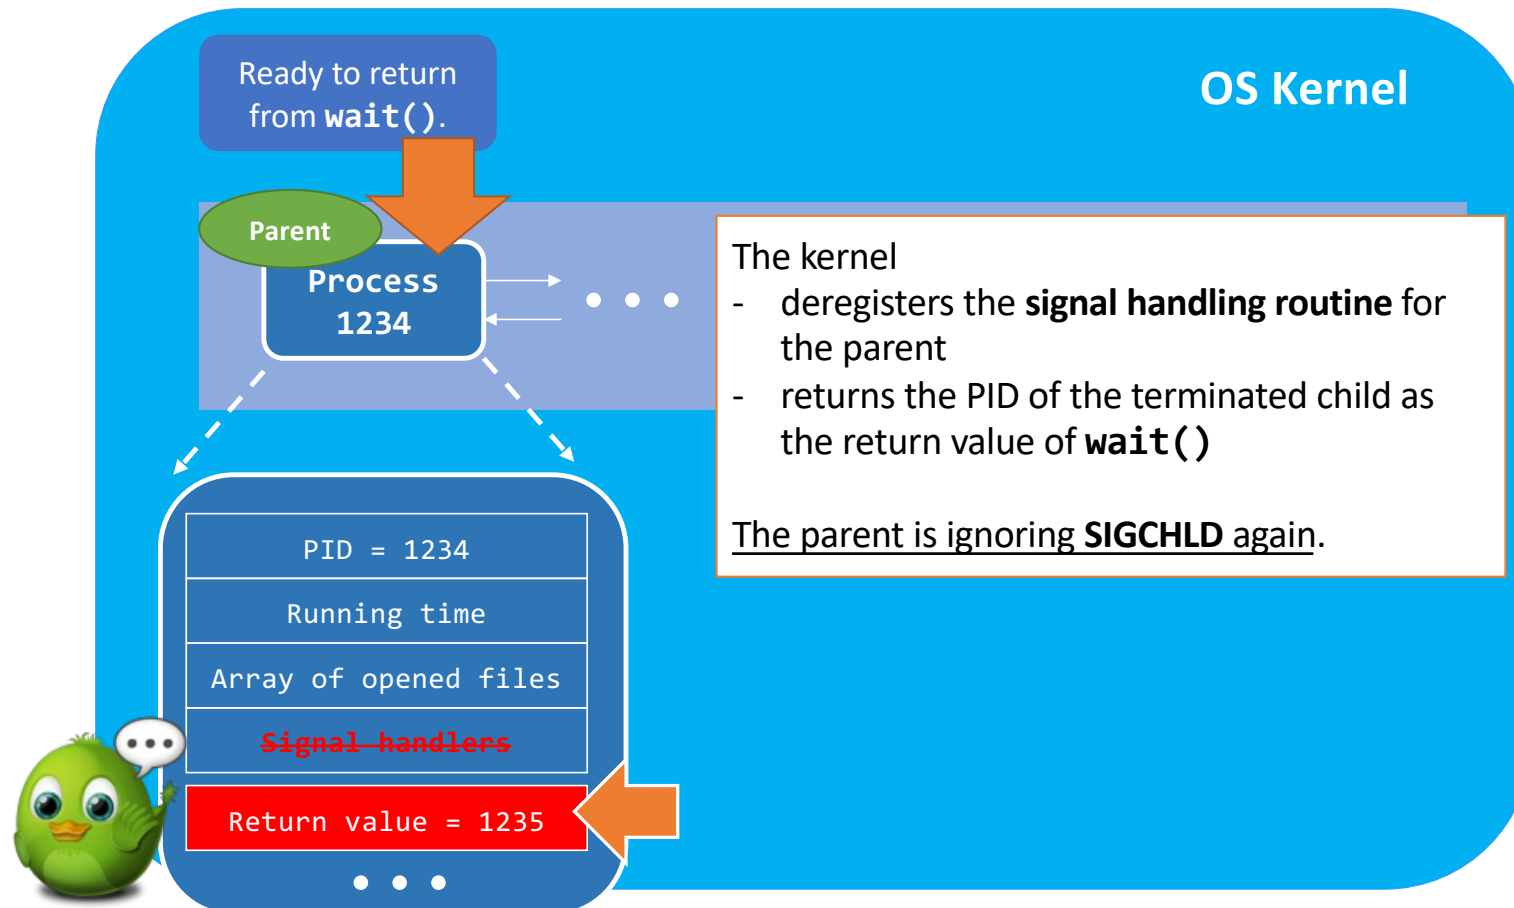
wait() Kernel View



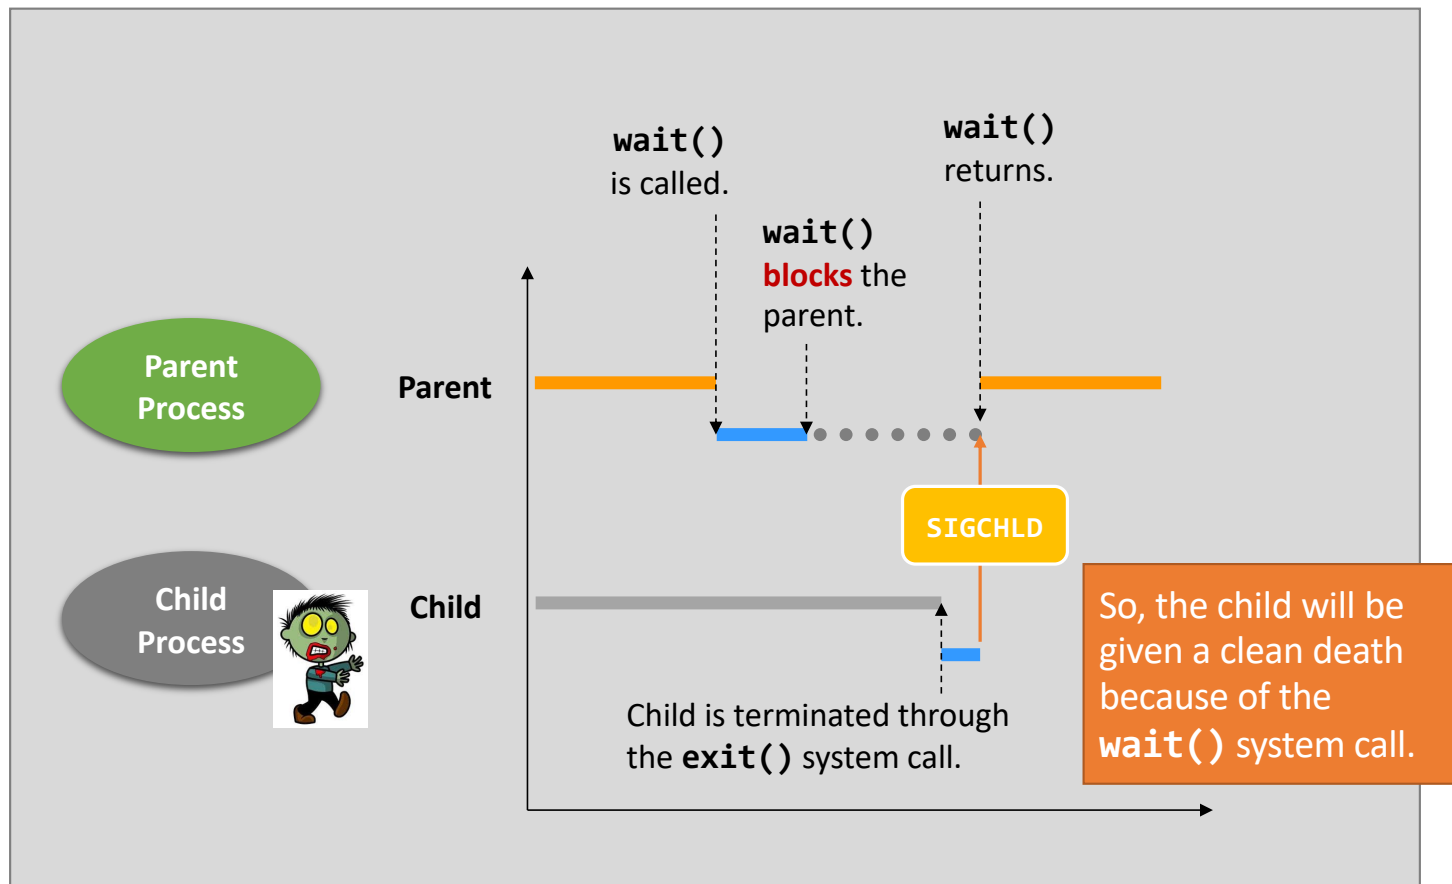
wait() Kernel View



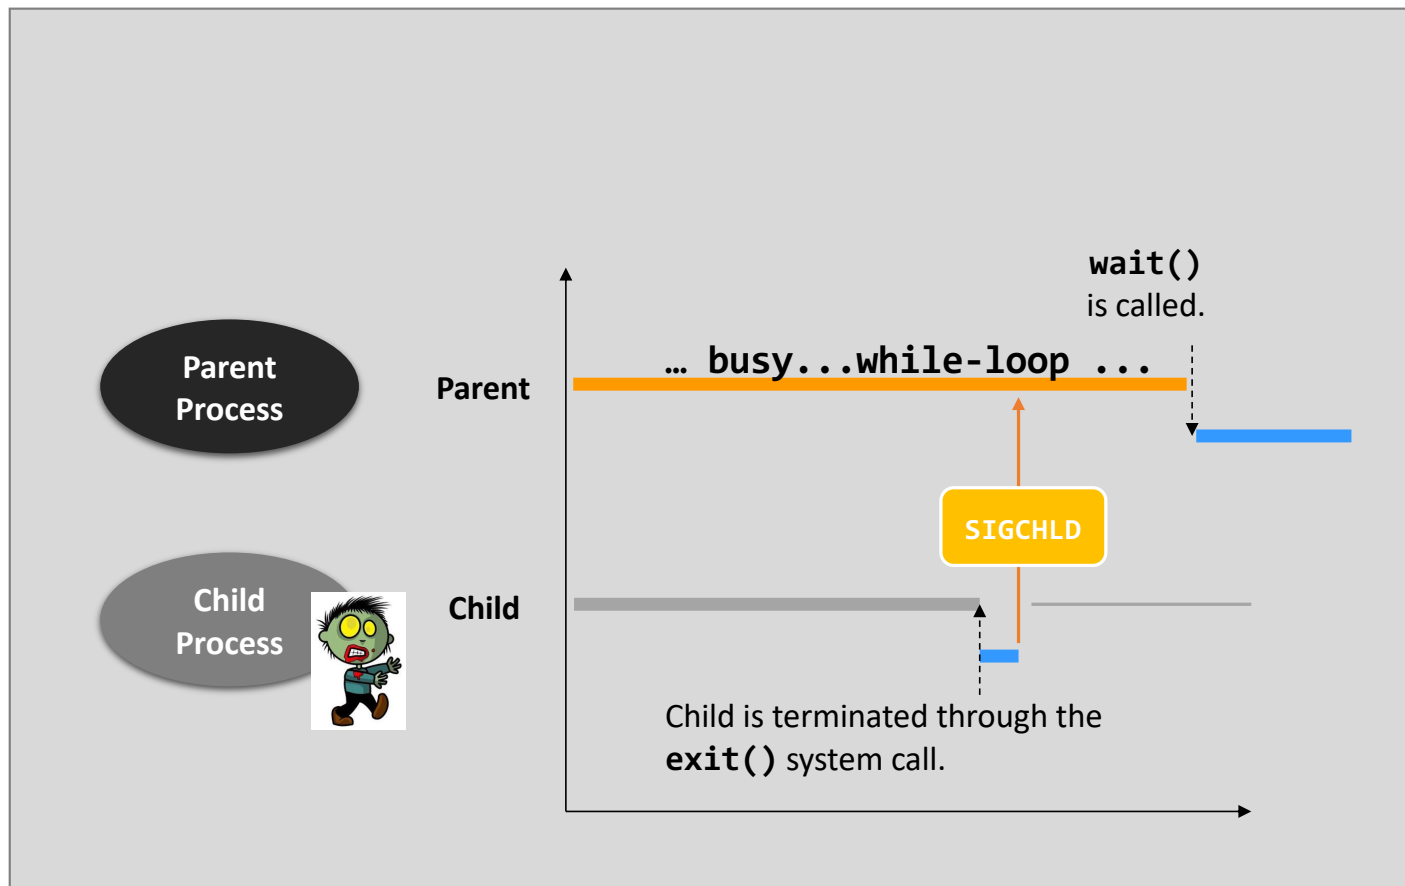
wait() Kernel View



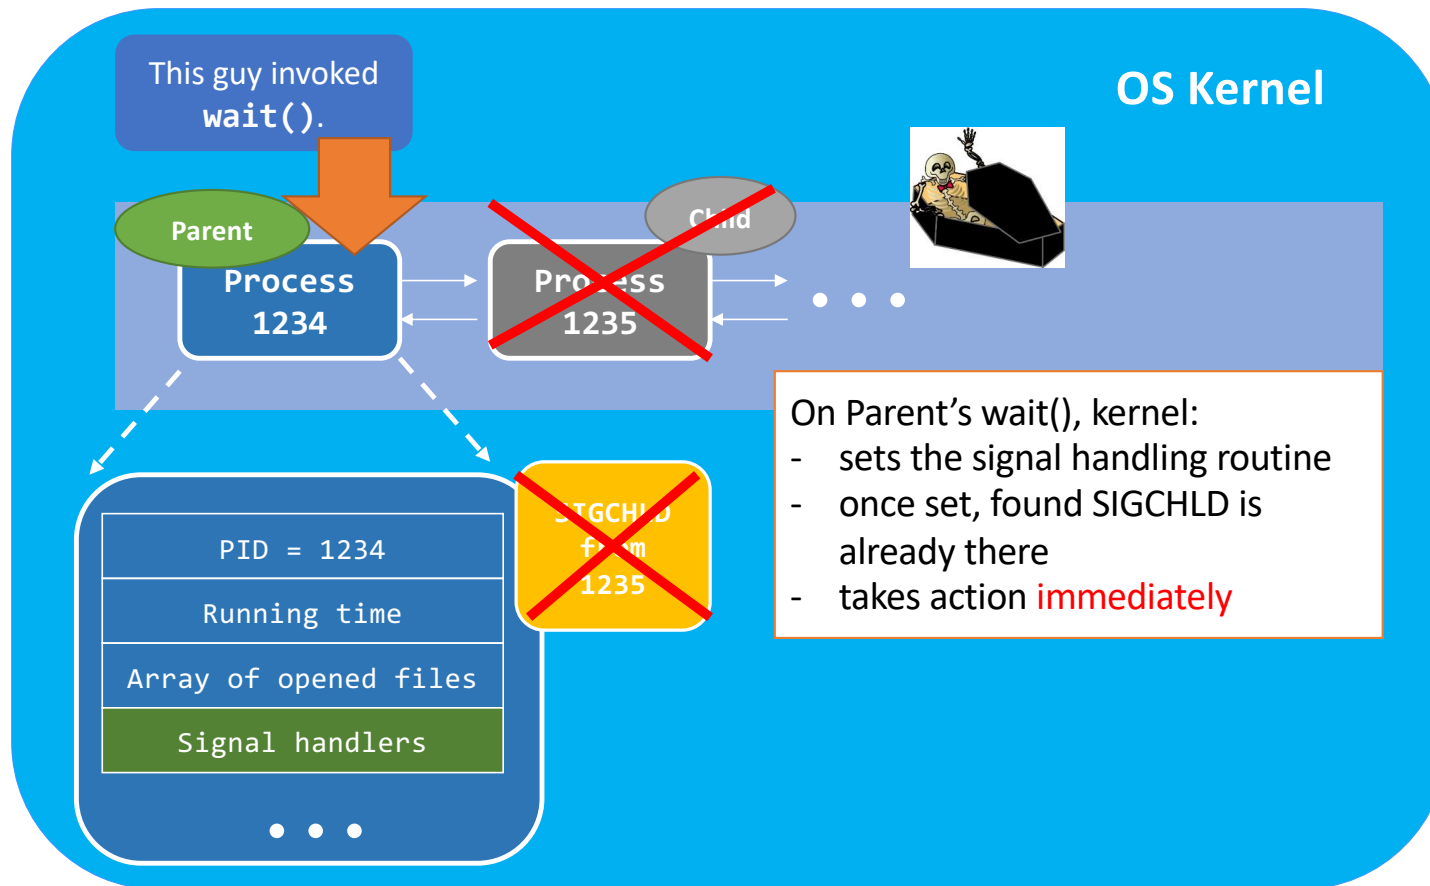
Normal Case



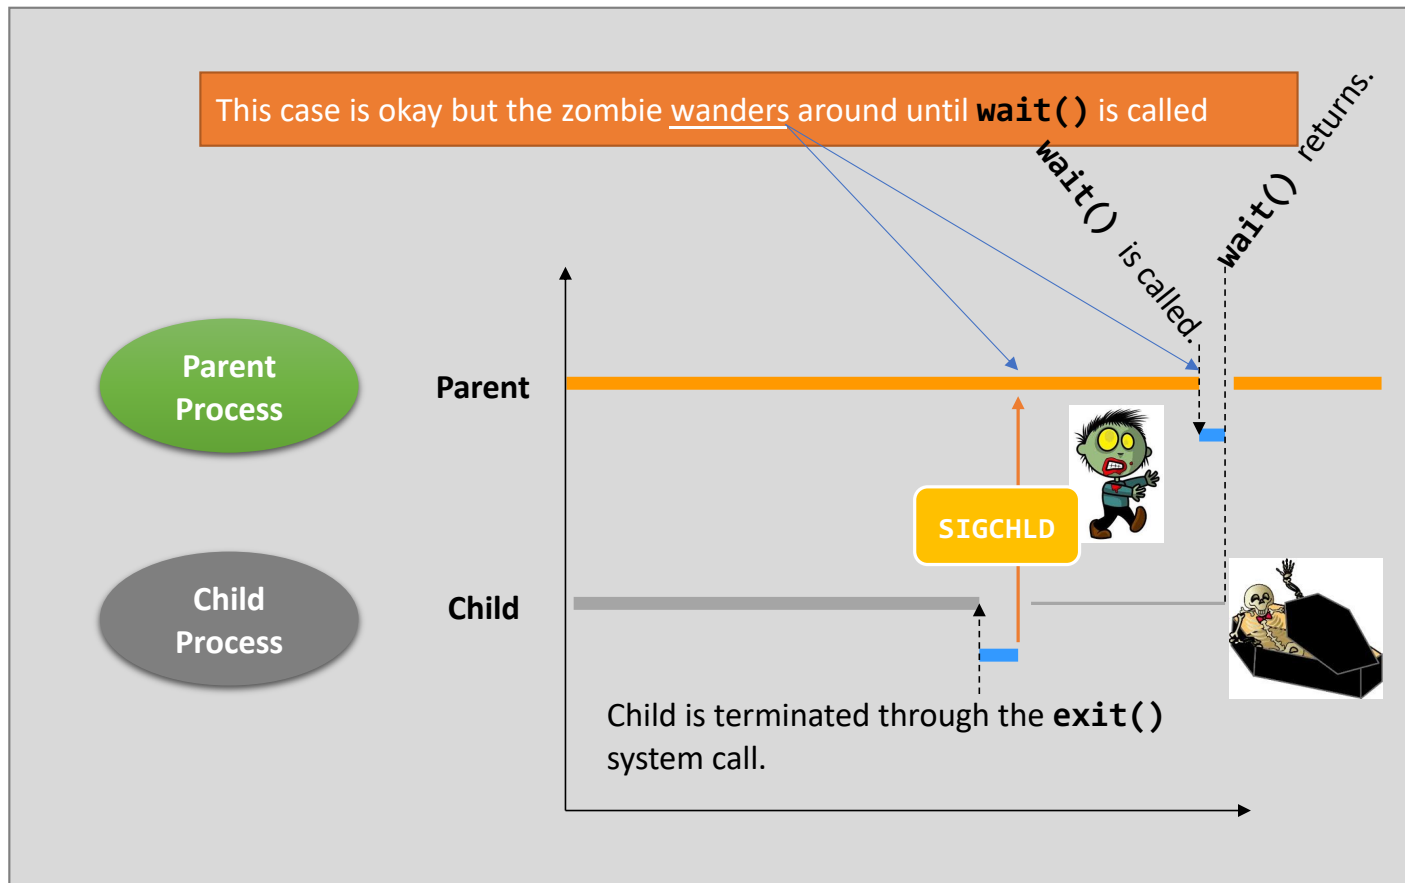
Parent's wait() after Child's exit()



Parent's Wait() after Child's exit()



Parent's wait() after Child's exit()



Summary of `wait()` and `exit()`

- `exit()` system call turns a process into a zombie when...
 - The process calls `exit()`.
 - The process returns from `main()`.
 - The process terminates abnormally.
 - The kernel knows that the process is terminated abnormally. Hence, the kernel invokes `exit()` for it.

Summary of wait() and exit()

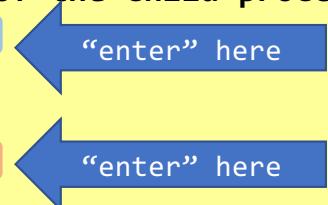
- `wait()` & `waitpid()` reap zombie child processes.
 - It is a must that you should never leave any zombies in the system.
 - `wait()` & `waitpid()` pause the caller until
 - A child terminates/stops, OR
 - The caller receives a signal (i.e., the signal interrupted the `wait()`)
- Linux will label zombie processes as “<defunct>”.
 - To look for them:

```
$ ps aux | grep defunct
..... 3150 ... [ls] <defunct>
$ _
```

PID of the
process

Summary of wait() and exit()

```
1 int main(void)
2 {
3     int pid;
4     if( (pid = fork()) !=0 ) {
5         printf("Look at the status of the child process %d\n", pid);
6         while( getchar() != '\n' );
7         wait(NULL);
8         printf("Look again!\n");
9         while( getchar() != '\n' );
10    }
11    return 0;
12 }
```



This program requires you to type “enter” twice before the process terminates.



You are expected to see **the status of the child process changes (ps aux [PID])** between the 1st and the 2nd “enter”.

Using wait() for Resource Management

- It is not only about process execution / suspension...
- It is about system resource management.
 - A zombie takes up a PID;
 - The total number of PIDs are limited;
 - Read the limit: "cat /proc/sys/kernel/pid_max"
 - It is 32,768.
 - What will happen if we don't clean up the zombies?

Using wait() for Resource Management

```
int main(void) {  
    while( fork() );  
    return 0;  
}
```



```
$ ./interesting
```

```
-
```

Terminal A

```
$ ls  
No process left.  
$ poweroff  
No process left.  
$ ==  
No process left.  
$ -
```

Terminal B

More about Processes

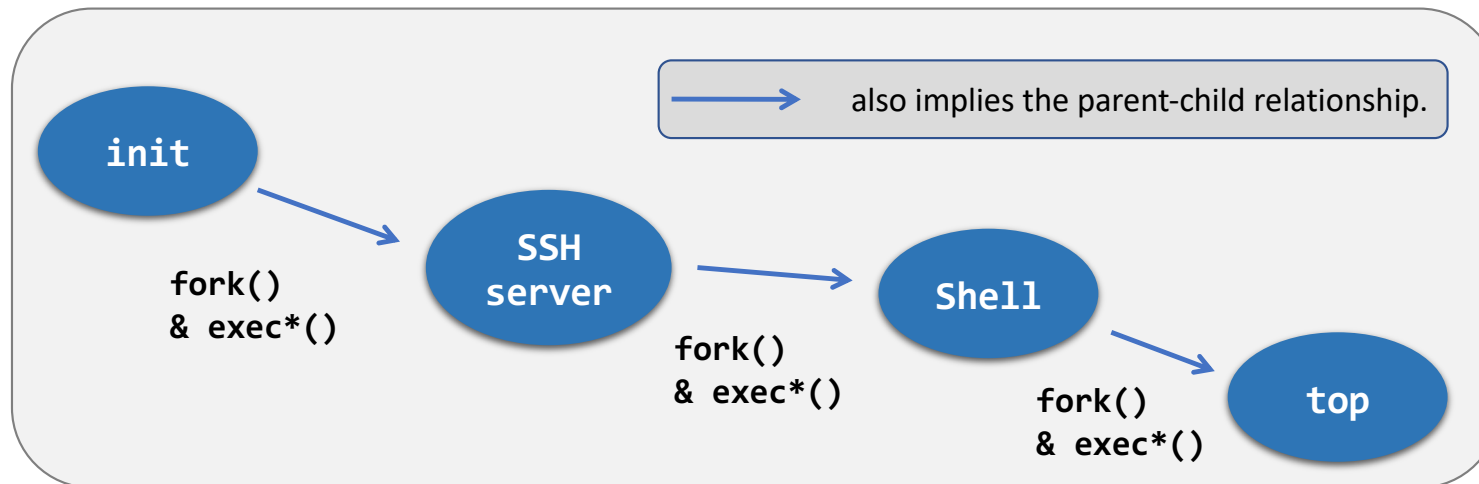
The first process

- We now focus on the process-related events.
 - The kernel, while it is booting up, creates the first process – init.
- The “init” process:
 - has PID = 1, and
 - is running the program code “/sbin/init”.
- Its first task is to create more processes...
 - Using fork() and exec().

How does uCore
create the first
process?


A Tree of Processes

- You can view the tree with the command:
 - “pstree”; or
 - “pstree -A” for ASCII-character-only display.



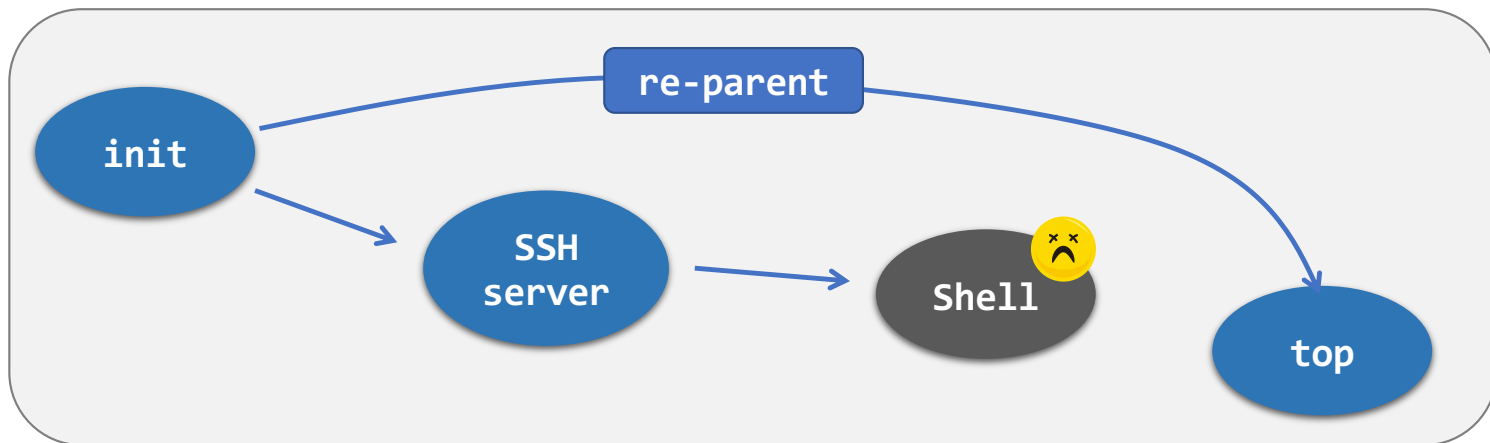
Orphans

- However, termination can happen, at any time and in any place...
 - This is no good because an orphan turns the hierarchy from a **tree** into a **forest**!
 - Plus, no one would know the termination of the orphan.



Re-parent

- In Linux
 - The “init” process will become the step-mother of all orphans
 - It's called **re-parenting**
- In Windows
 - It maintains a *forest-like process hierarchy*.....



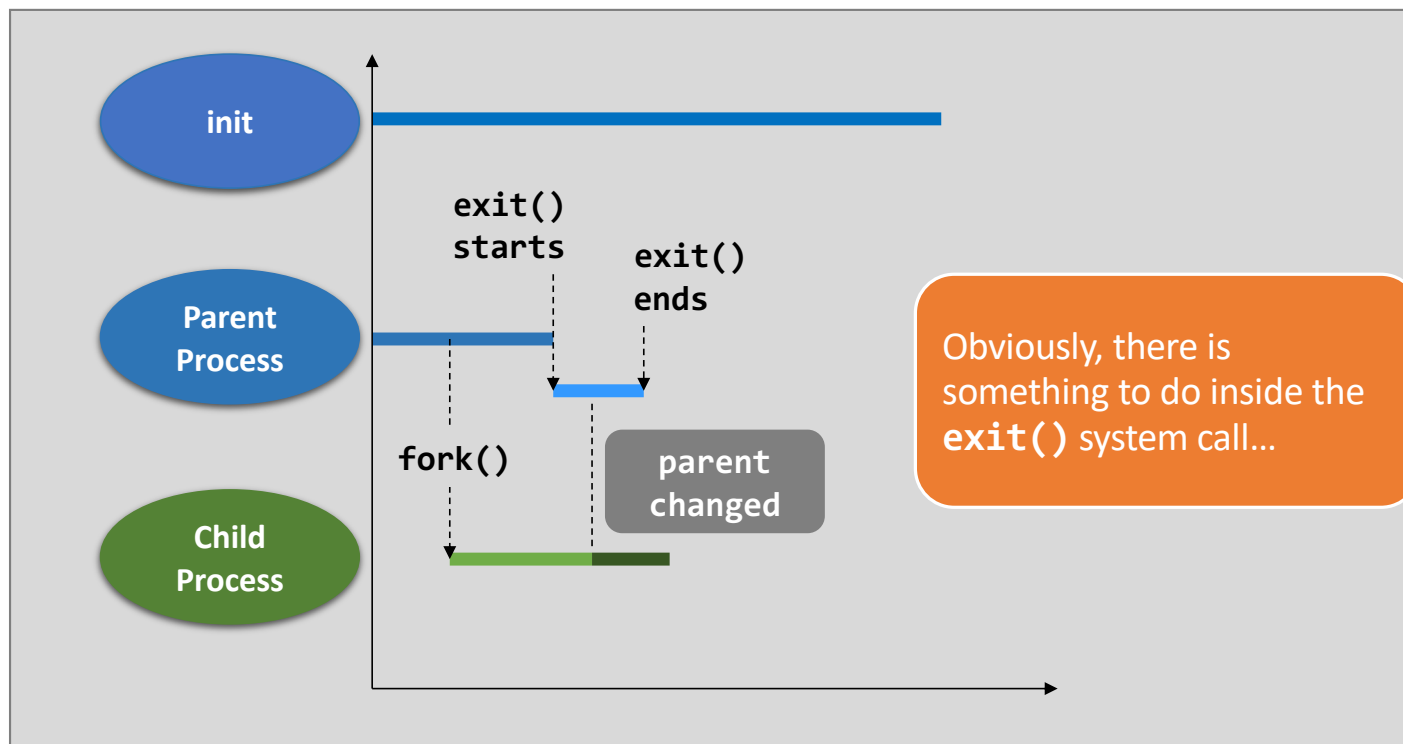
An Example

```
1 int main(void) {
2     int i;
3     if(fork() == 0) {
4         for(i = 0; i < 5; i++) {
5             printf("(%d) parent's PID = %d\n",
6                 getpid(), getppid() );
7             sleep(1);
8         }
9     }
10    else
11        sleep(1);
12    printf("(%d) bye.\n", getpid());
13 }
```


getppid() is the system call that returns the parent's PID of the calling process.

```
$ ./reparent
(1235) parent's PID = 1234
(1235) parent's PID = 1234
(1234) bye.
$ (1235) parent's PID = 1
(1235) parent's PID = 1
(1235) parent's PID = 1
(1235) bye.
$ _
```

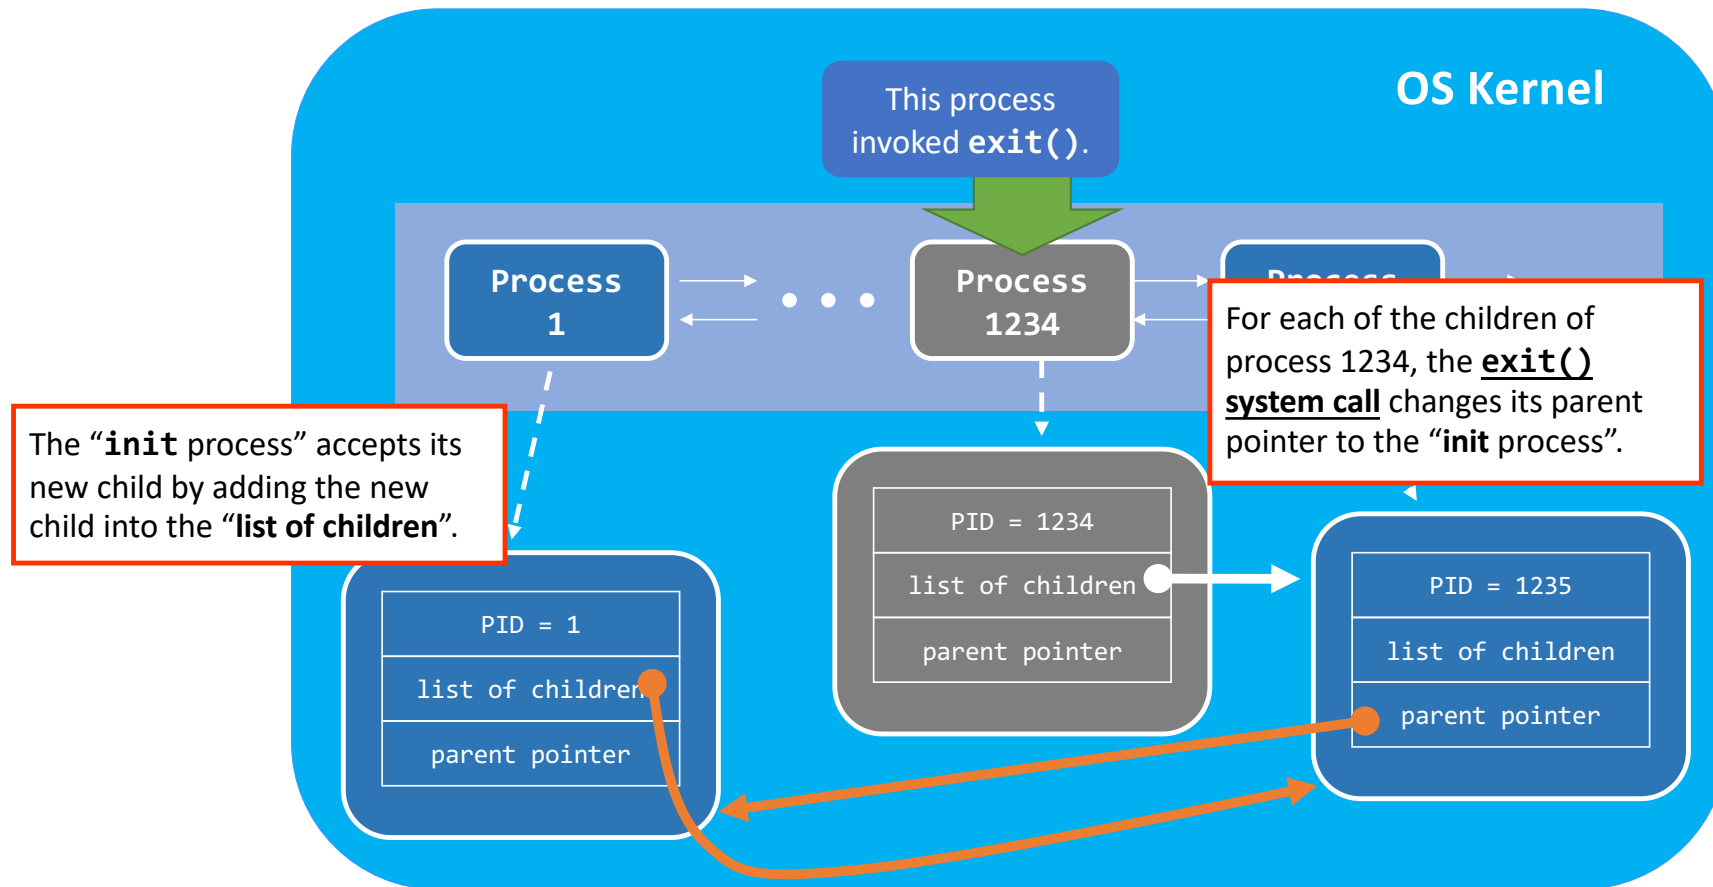
Re-parenting Explained



Re-parenting Explained (Cont'd)



Re-parenting Explained (Cont'd)



Background Jobs

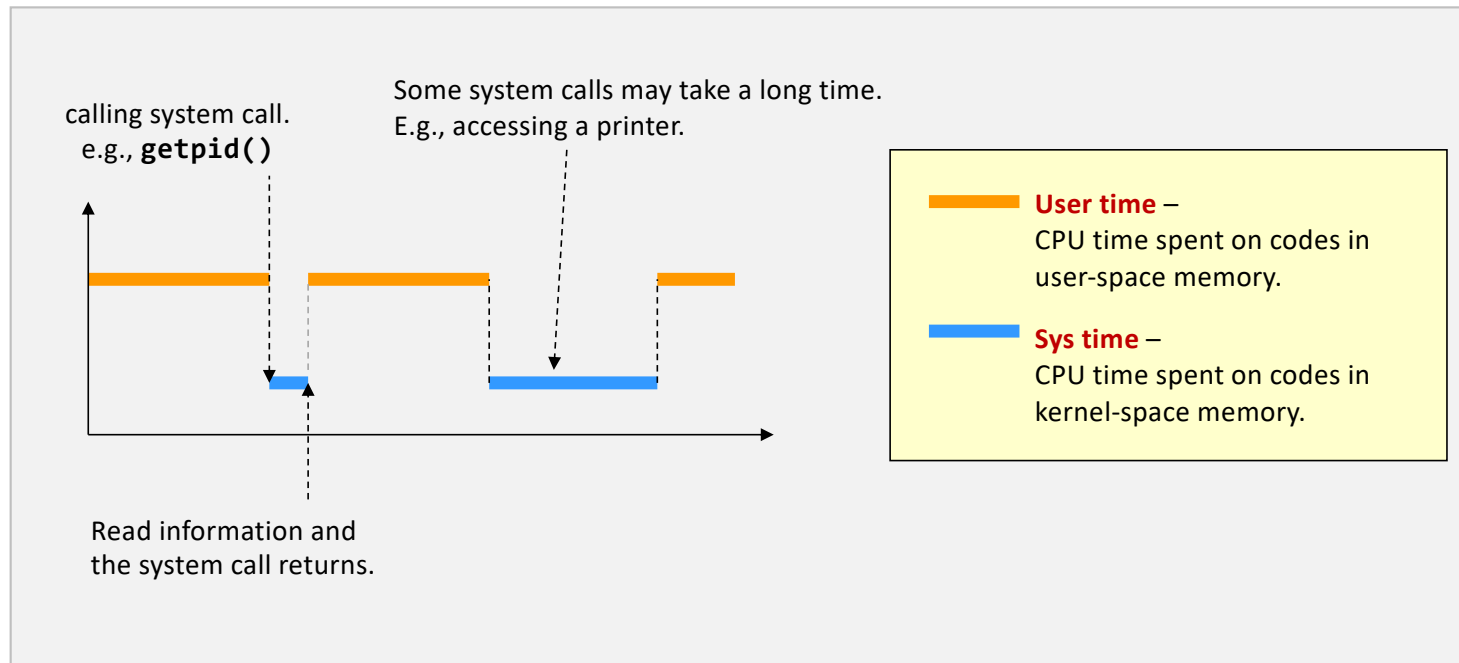
- The re-parenting operation enables something called **background jobs** in Linux
 - It allows a process runs **without a parent terminal/shell**

[Back to home](#)

```
$ ./infinite_loop &  
$ exit  
  
[ The shell is gone ]
```

```
$ ps -C infinite_loop  
PID  TTY  
1234  ... ./infinite_loop  
$ _
```

Measure Process Time



User Time v.s. System Time (Case 1)

```
$ time ./time_example
```

```
real    0m0.001s
user    0m0.000s
sys     0m0.000s
$ _
```

Real-time elapsed when “./time_example” terminates.

The user time of “./time_example”.

The sys time of “./time_example”.

It's possible:
real > user + sys
real < user + sys

Why?

- real > user + sys
 I/O intensive
- real < user + sys
 multi-core

```
int main(void) {
    int x = 0;
    for(i = 1; i <= 10000; i++) {
        x = x + i;
        // printf("x = %d\n", x);
    }
    return 0;
}
```

User Time v.s. System Time (Case 1)

```
$ time ./time_example
```

```
real    0m0.001s
user    0m0.000s
sys     0m0.000s
$ _
```

```
int main(void) {
    int x = 0;
    for(i = 1; i <= 10000; i++) {
        x = x + i;
        // printf("x = %d\n", x);
    }
    return 0;
}
```

Commented on purpose.

```
$ time ./time_example
```

```
real 0m2.795s
user 0m0.084s
sys 0m0.124s
$ _
```

See? Accessing hardware costs the process more time.

```
int main(void) {
    int x = 0;
    for(i = 1; i <= 10000; i++) {
        x = x + i;
        printf("x = %d\n", x);
    }
    return 0;
}
```


User Time v.s. System Time (Case 2)

- The user time and the sys time together **define the performance of an application.**
 - When writing a program, you must consider both the user time and the sys time.
 - E.g., the output of the following two programs are exactly the same. But, their running time is not.

```
#define MAX 1000000

int main(void) {
    int i;
    for(i = 0; i < MAX; i++)
        printf("x\n");
    return 0;
}
```

```
#define MAX 1000000

int main(void) {
    int i;
    for(i = 0; i < MAX / 5 ; i++)
        printf("x\nx\nx\nx\nx\n");
    return 0;
}
```

User Time v.s. System Time (Case 2)

```
#define MAX 1000000

int main(void) {
    int i;
    for(i = 0; i < MAX; i++)
        printf("x\n");
    return 0;
}
```

```
$ time ./time_example_slow

real 0m1.562s
user 0m0.024s
sys  0m0.108s
$ _
```

```
#define MAX 1000000

int main(void) {
    int i;
    for(i = 0; i < MAX / 5 ; i++)
        printf("x\nx\nx\nx\nx\n");
    return 0;
}
```

```
$ time ./time_example_fast

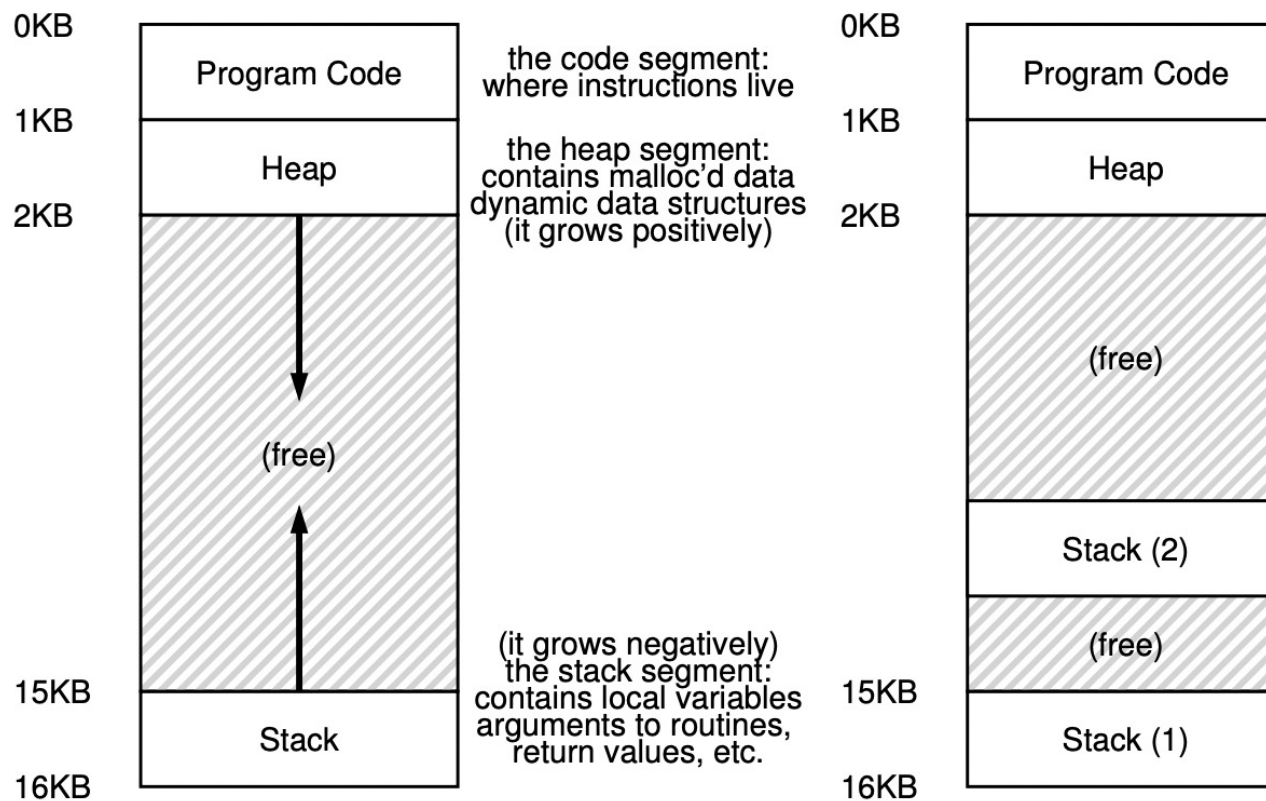
real 0m1.293s
user 0m0.012s
sys  0m0.084s
$ _
```

Threads

What is a Thread?

- Thread is an **abstraction** of the execution of a program
 - A single-threaded program has one point of execution
 - A multi-threaded program has more than one points of execution
- Each thread has its own **private** execution state
 - Program counter and a private set of registers
 - A private stack for thread-local storage
 - CPU switching from one thread to another requires context switch
- Threads in the same process **share** computing resources
 - Address space, files, signals, etc.

Single-Threaded and Multi-Threaded



Why Use Thread?


- Increase parallelism
 - One thread per CPU makes better use of multiple CPUs to improve efficiency
- Avoid blocking program progress due to slow I/O
 - Threading enables overlap of I/O with other activities within a single program
 - e.g., many modern server-based applications (web servers, database management systems, and the like) make use of threads
- And allow resource sharing !!!

Thread Implementation


- User-level thread
 - Thread management (e.g., creating, scheduling, termination) done by user-level threads library
 - OS does not know about user-level thread
- Kernel-level thread
 - Thread management done by kernel
 - OS is aware of each kernel-level thread

Thread Models


- One-to-one mapping
 - One user-level thread to one kernel-level thread
- Many-to-one mapping
 - Many user-level thread to one kernel-level thread
- Many-to-many mapping
 - Many user-level thread to many kernel-level thread



One-to-One



Many-to-One



Many-to-Many

Pros and Cons

- Many-to-one mapping
 - Pros: context switch between threads is cheap
 - Cons: When one thread blocks on I/O, all threads block
- One-to-one mapping
 - Pros: Every thread can run or block independently
 - Cons: Need to make a crossing into kernel mode to schedule
- Many-to-many mapping
 - Many user-level threads multiplexed on less or equal number of kernel-level threads
 - Pros: best of the two worlds, more flexible
 - Cons: difficult to implement

Thank you!

